

CPSC 340:


Machine Learning and Data Mining

Introduction to Reinforcement Learning -- Bonus Lecture

Helen Zhang
(slides adapted from Daniele Reda)
Fall 2021

Today's Plan:

- What is RL
- Funny videos
- Q-learning, DQN
- Self-driving car



Law of Effect

"responses that produce a satisfying effect in a particular situation become more likely to occur again in that situation, and responses that produce a discomforting effect become less likely to occur again in that situation."

Edward Thorndike



Positive

Negative

REWARD


REWARD

$$Q(s, a) = Q(s, a) + \alpha[R(s, a) + \gamma \max_{a'} Q'(s', a') - Q(s, a)]$$


What is Reinforcement Learning?

- learning by trial and error
- learning by interacting with environment


Problem Setting



Alpha Go



Branches of Machine Learning



Supervised Learning

Data	Label
X1	Y1
X2	Y2
X3	Y3
...	...

Supervised Learning




Representation of the data



Output: CAT

Unsupervised Learning




Characteristics

- no supervisor, only a reward signal
- feedback is delayed, not instantaneous
- process is iterative (time matters)
- agent's actions affect subsequent data it receives


Flipping pancakes



Walking simulation



Problem Setting



Goal

- maximize total reward

$$R_t = r_{t+1} + \gamma r_{t+2} + \gamma^2 r_{t+3} + \cdots + \gamma^{T-t-1} r_T$$


- T to infinity
- Why discount factor?
 - convergence
 - sooner rewards are usually more useful than later ones

Value Function

$$V^*(S) = \max_a [R(s, a) + \gamma \sum_{s'} p(s'|s, a) V^*(s')]$$

- Problems:
 - it doesn't tell which actions to take

Representation of the system




Q-Value Function

$$Q^*(s, a) = R(s, a) + \gamma \sum_{s'} p(s' \mid s, a) \max_{\alpha} (Q^*(s', \alpha))$$

A very simple algorithm: Q-learning

		Actions					
States		0	1	2	3	4	5
	0	-1	-1	-1	-1	0	-1
	1	-1	-1	-1	0	1	100
	2	-1	-1	-1	0	-1	-1
	3	-1	0	0	-1	0	-1
	4	0	-1	-1	0	-1	100
	5	-1	0	-1	-1	0	100

Exploration vs Exploitation



One strategy: ϵ -greedy

Q-learning

Initialize Q-table with random values.

1. Choose action a to perform in current state s . (ϵ -greedy)
2. Perform a and receive reward $R(s,a)$.
3. Observe new state $S(s,a)$.
4. Update Q-table.

$$Q'(s, a) \leftarrow R(s, a) + \gamma \max_{\alpha} \{Q'(\mathcal{S}(s, a), \alpha)\}$$

PROBLEM:
TABLE CAN EASILY EXPLODE IN DIMENSIONS

Let's look at an example: ATARI



State (the actual image)
84x84x4 pixels (gray-scale)



2 Actions
Left-Right

Could we use a q-table?

Atari Breakout example:

State = raw pixels of last 4 frames (84x84) with 256 different possible values.

Actions = 2 actions available




$(4 \times 84 \times 84 \times 256) \times 2 = 14450688$ different values

Solution

Neural networks!

Neural Networks



Neural Networks

$$loss = \left(\underbrace{r + \gamma \max_{a'} \hat{Q}(s, a')}_{\text{Target}} - \underbrace{Q(s, a)}_{\text{Prediction}} \right)^2$$

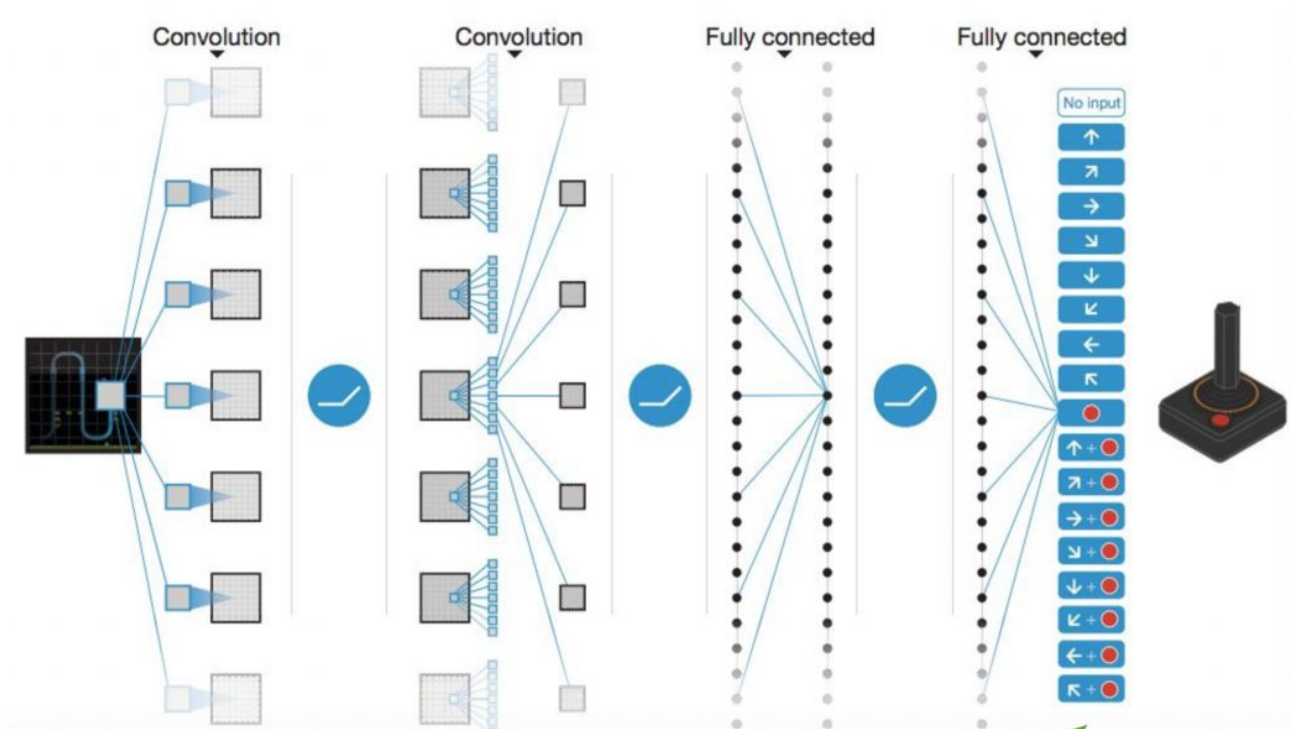
Diagram illustrating the loss function for a neural network in a reinforcement learning context. The loss is calculated as the squared difference between the Target and the Prediction.

The Target is composed of the Reward (r) and the discounted maximum future value ($\gamma \max_{a'} \hat{Q}(s, a')$).

The Prediction is the current value function ($Q(s, a)$).

$$Q'(s, a) \leftarrow \mathcal{R}(s, a) + \gamma \max_{\alpha} \{Q'(\mathcal{S}(s, a), \alpha)\}$$

DQN Framework



1 network, outputs Q value for each action


DQN Framework

Algorithm 1: DQN Pesudocode

```
1 Randomly initialize neural network  $NN$ 
2 Get initial state  $s_0$ 
3  $t = 0$ 
4 while 1 do
    //  $\epsilon$ -greedy strategy
5      $r = \text{get random value from } (0, 1)$ 
6     if  $r > \epsilon$  then
7          $a_t = NN(s_0)$ 
8     else
9          $a_t = \text{random action between the ones available}$ 
10    end
11     $s_{t+1}, r_t, done = \text{environment}(a_t)$ 
12    if  $done = \text{True}$  then
13         $y = r_t$ 
14    else
15         $y = r + \gamma * \max_{a_i} NN(s_{t+1}$ 
16    end
17    Do gradient descent on  $y - NN(s_t, a_t)$  to update weights of  $NN$ 
18     $t = t + 1$ 
19 end
```

Famous successes of RL: Atari Breakout






A real world example: Learning to drive with Reinforcement Learning



Learning to drive with RL



**Steering & Speed
Measurement**



**Steering & Speed
Command**

Reward: forward distance

Terminate when it goes out of the lane



Kendall, Alex, et al. "Learning to drive in a day."

Where to go from here?

- [openAI Spinning Up RL](#)
- [Sutton book](#)
- [David Silver UCL Course](#)

The right side of the slide features a vertical panel with an abstract design of thin, overlapping, colorful lines in shades of blue, green, yellow, and pink, creating a sense of movement and depth.

Reinforcement Learning

An Introduction
second edition

Richard S. Sutton and Andrew G. Barto