

CPSC 340: Machine Learning and Data Mining

Sparse Matrix Factorization

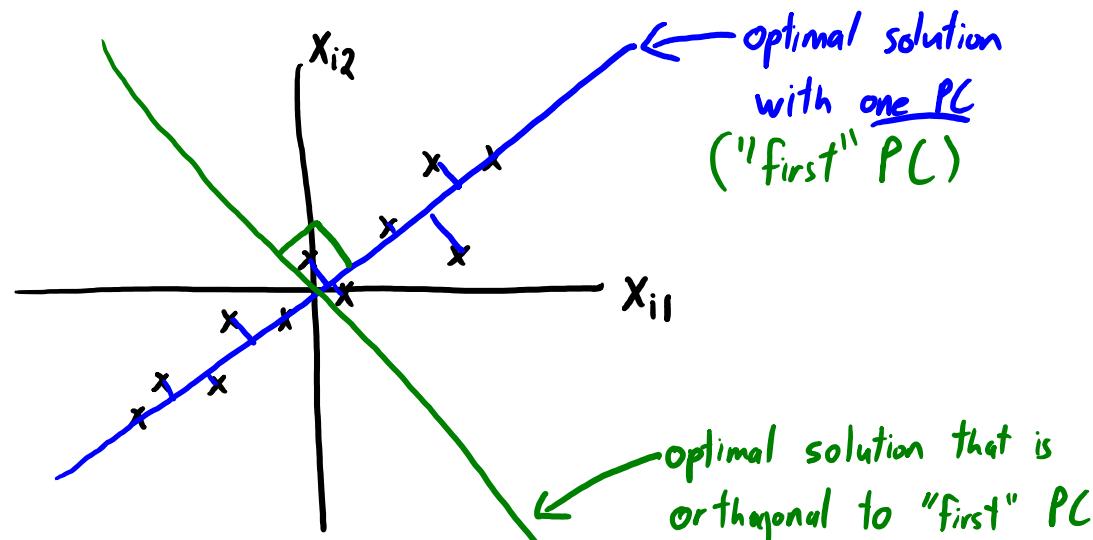
Fall 2021

Last Time: PCA with Orthogonal/Sequential Basis

- When $k = 1$, PCA has a **scaling problem**.
- When $k > 1$, have **scaling, rotation, and label switching**.
 - Standard fix: use **normalized orthogonal rows W_c** of ' W '.

$$\|w_c\|=1 \quad \text{and} \quad w_c^T w_{c'} = 0 \quad \text{for } c' \neq c$$

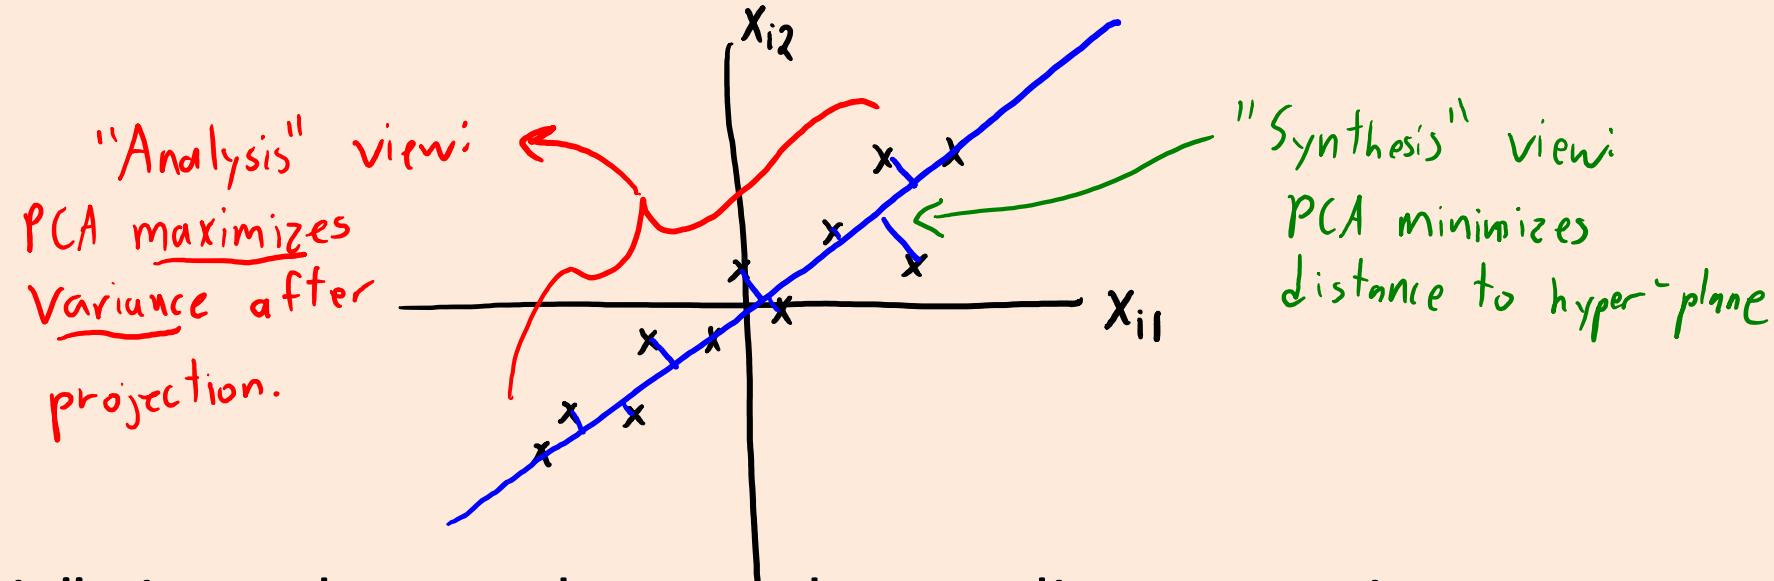
- And **fit the rows in order**:
 - First row “explains the most variance” or “reduces error the most”.



bonus!

“Synthesis” View vs. “Analysis” View

- We said that PCA finds hyper-plane minimizing distance to data x_i .
 - This is the “synthesis” view of PCA (connects to k-means and least squares).



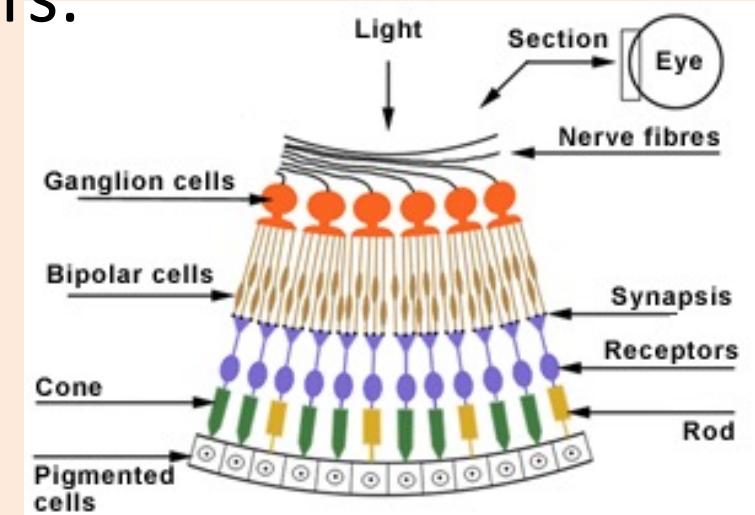
- “Analysis” view when we have orthogonality constraints:
 - PCA finds hyper-plane maximizing variance in z_i space.
 - You pick W to “explain as much variance in the data” as possible.

bonus!

Colour Opponency in the Human Eye

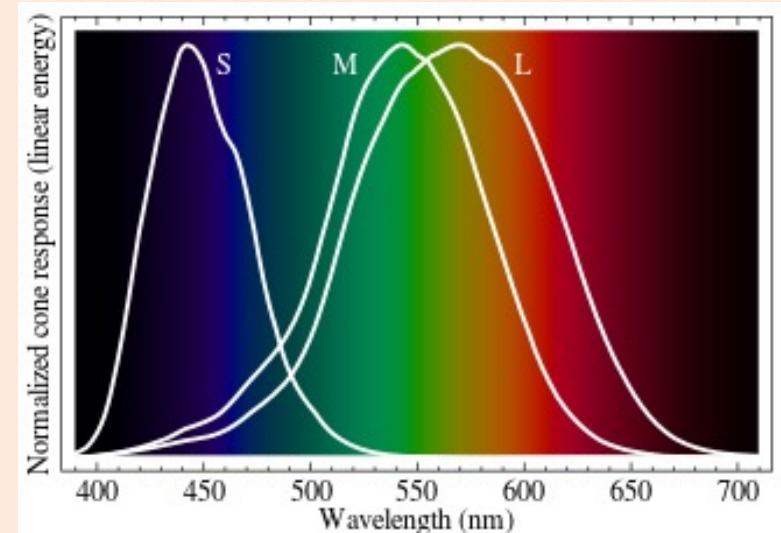
- Classic model of the eye is with 4 photoreceptors:

- Rods (more sensitive to brightness).
- L-Cones (most sensitive to red).
- M-Cones (most sensitive to green).
- S-Cones (most sensitive to blue).



- Two problems with this system:

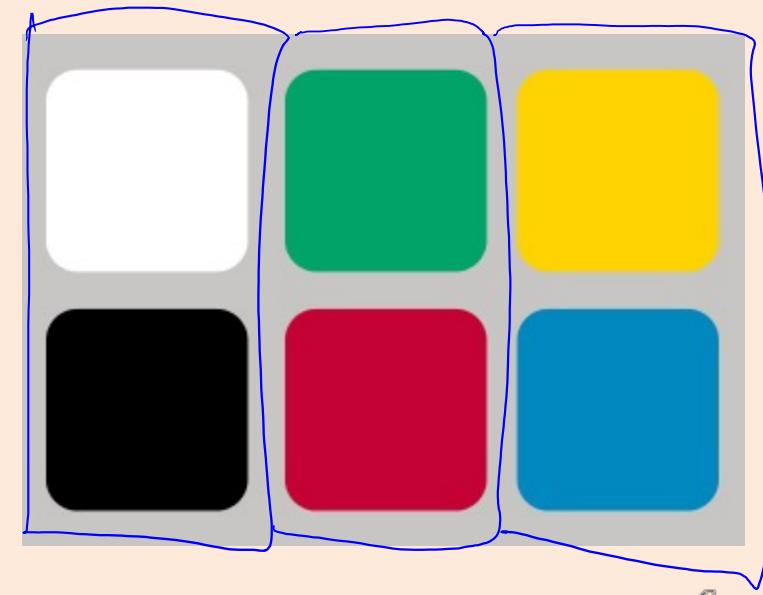
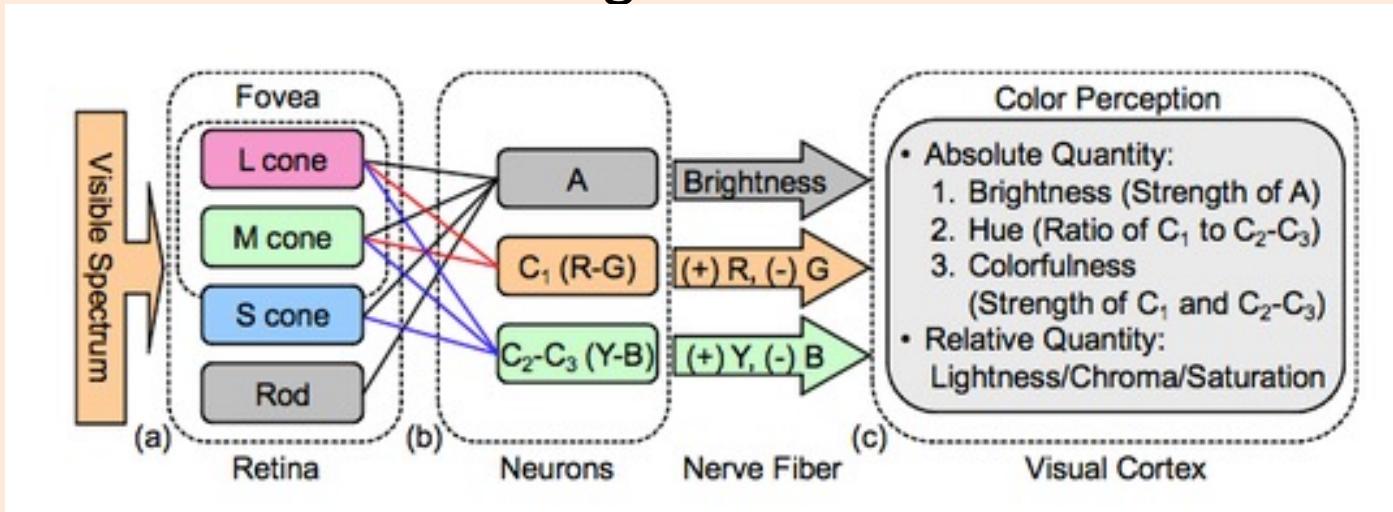
- Not orthogonal.
 - High correlation in particular between red/green.
- We have 4 receptors for 3 colours.



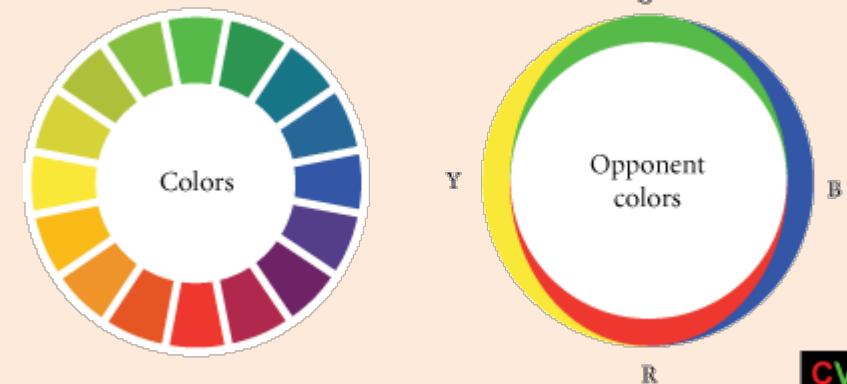
bonus!

Colour Opponency in the Human Eye

- Bipolar and ganglion cells seem to code using “opponent colors”:
 - 3-variable orthogonal basis:

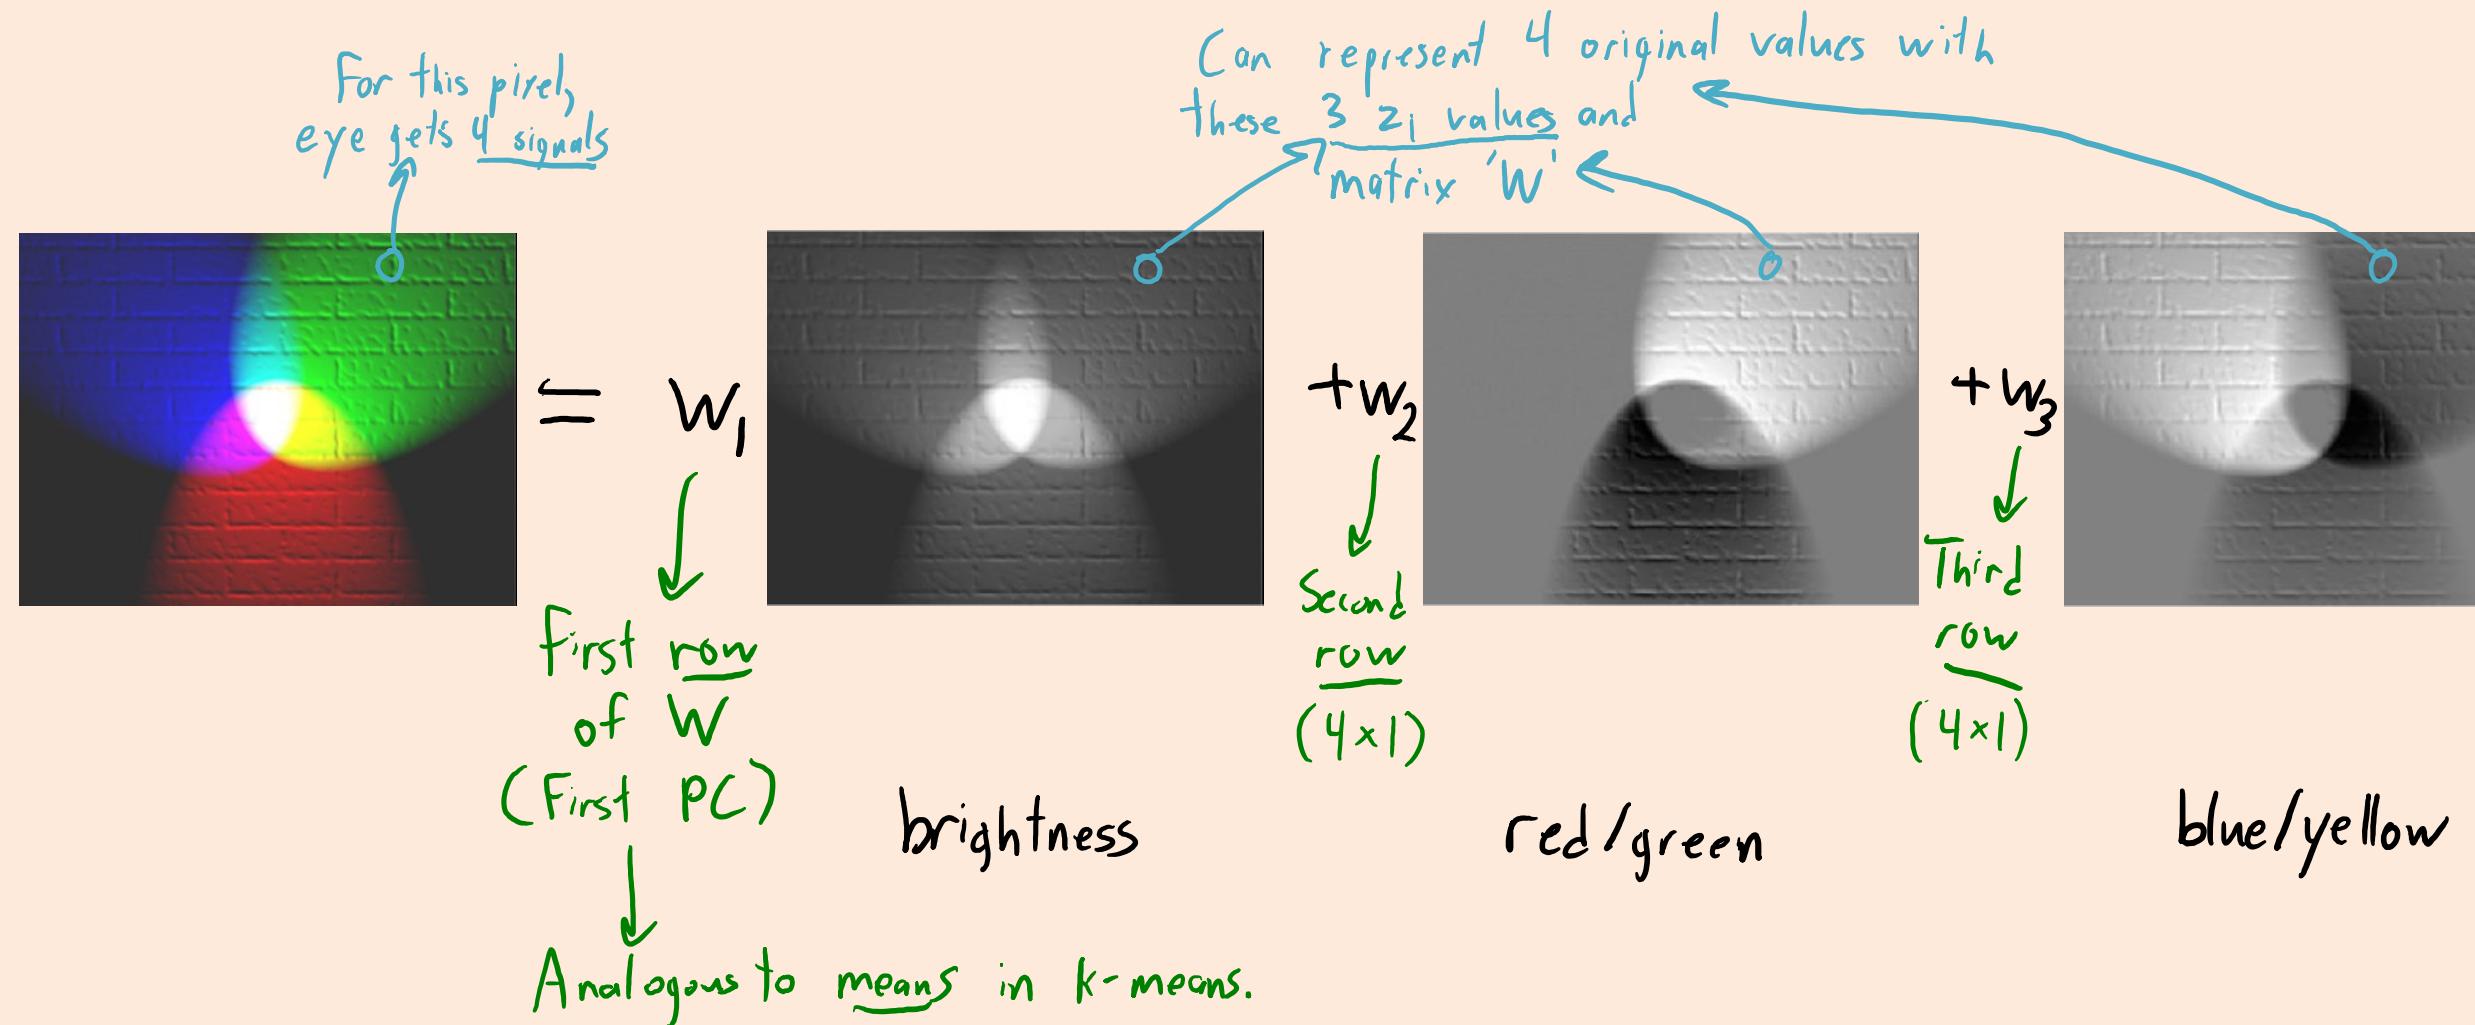


- This is similar to PCA ($d = 4$, $k = 3$).



bonus!

Colour Opponency Representation



PCA Computation: other methods

- With linear regression, we had the normal equations
 - But we also could do it with gradient descent, SGD, etc.
- With PCA we have the SVD
 - But we can also do it with gradient descent, SGD, etc.
 - These other methods typically don't enforce the uniqueness "constraints".
 - Sensitive to initialization, don't enforce normalization, orthogonality, ordered PCs.
 - But you can do this in post-processing if you want.
 - Why would we want this? We can use our tricks from Part 3 of the course:
 - We can do things like "robust" PCA, "regularized" PCA, "sparse" PCA, "binary" PCA.
 - We can fit huge datasets where SVD is too expensive.

PCA Computation: Alternating Minimization

- With centered data, the PCA objective is:

$$f(W, Z) = \sum_{i=1}^n \sum_{j=1}^d (\langle w_j^i, z_i \rangle - x_{ij})^2$$

- In k-means we tried to optimize this with alternating minimization:
 - Fix “cluster assignments” Z and find the optimal “means” W.
 - Fix “means” W and find the optimal “cluster assignments” Z.
- Converges to a local optimum.
 - But may not find a global optimum (sensitive to initialization).

PCA Computation: Alternating Minimization

- With centered data, the PCA objective is:

$$f(W, Z) = \sum_{i=1}^n \sum_{j=1}^d (\langle w_i^T, z_i \rangle - x_{ij})^2$$

- In PCA we can also use alternating minimization:
 - Fix “part weights” Z and find the optimal “parts” W.
 - Fix “parts” W and find the optimal “part weights” Z.
- Converges to a local optimum.
 - Which will be a global optimum (if we randomly initialize W and Z).

PCA Computation: Alternating Minimization

- With centered data, the PCA objective is:

$$f(W, Z) = \sum_{i=1}^n \sum_{j=1}^d (\langle w_j^T, z_i \rangle - x_{ij})^2$$

- Alternating minimization steps:

- If we fix Z , this is a quadratic function of W (least squares column-wise):

$$\nabla_W f(W, Z) = Z^T Z W - Z^T X \quad \text{so} \quad W = \underbrace{(Z^T Z)^{-1}}_{\text{(writing gradient as a matrix)}} (Z^T X)$$

- If we fix W , this is a quadratic function of Z (transpose due to dimensions):

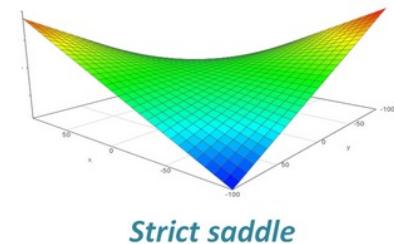
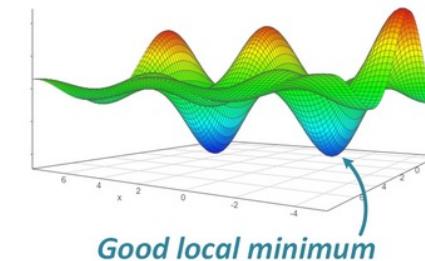
$$\nabla_Z f(W, Z) = Z W W^T - X W^T \quad \text{so} \quad Z = \underbrace{X W^T (W W^T)^{-1}}_T$$

These are usually invertible since $k < n$ and $k < d$

PCA Computation: Alternating Minimization

- With centered data, the PCA objective is:

$$f(W, Z) = \sum_{i=1}^n \sum_{j=1}^d (\langle w_i^T z_i \rangle - x_{ij})^2$$



- This objective is **not jointly convex** in W and Z .
 - You will **find different W and Z depending on the initialization.**
 - For example, if you initialize with all $w_c = 0$, then they will stay at zero.
 - But it's possible to show that all "stable" local optima are global optima.
 - You will **converge to a global optimum in practice** if you **initialize randomly**.
 - Randomization means you don't start on one of the unstable non-global critical points.
 - E.g., sample each initial z_{ij} from a normal distribution.

PCA Computation: Stochastic Gradient

- For big X matrices, you can also use **stochastic gradient**:

$$f(W, Z) = \sum_{i=1}^n \sum_{j=1}^d (\langle w_j^i, z_i \rangle - x_{ij})^2 = \sum_{(i,j)} (\underbrace{\langle w_j^i, z_i \rangle - x_{ij}}_{f(w_j^i, z_i, x_{ij})})^2$$

On each iteration, pick a random example i and feature j .

$$\rightarrow \text{Set } w_j^i \text{ to } w_j^i - \alpha^t \nabla_{w_j^i} f(w^i, z_i, x_{ij})$$

$$\rightarrow \text{Set } z_i \text{ to } z_i - \alpha^t \nabla_{z_i} f(w^i, z_i, x_{ij})$$

- Other variables stay the same, cost per iteration is only $O(k)$.

(pause)

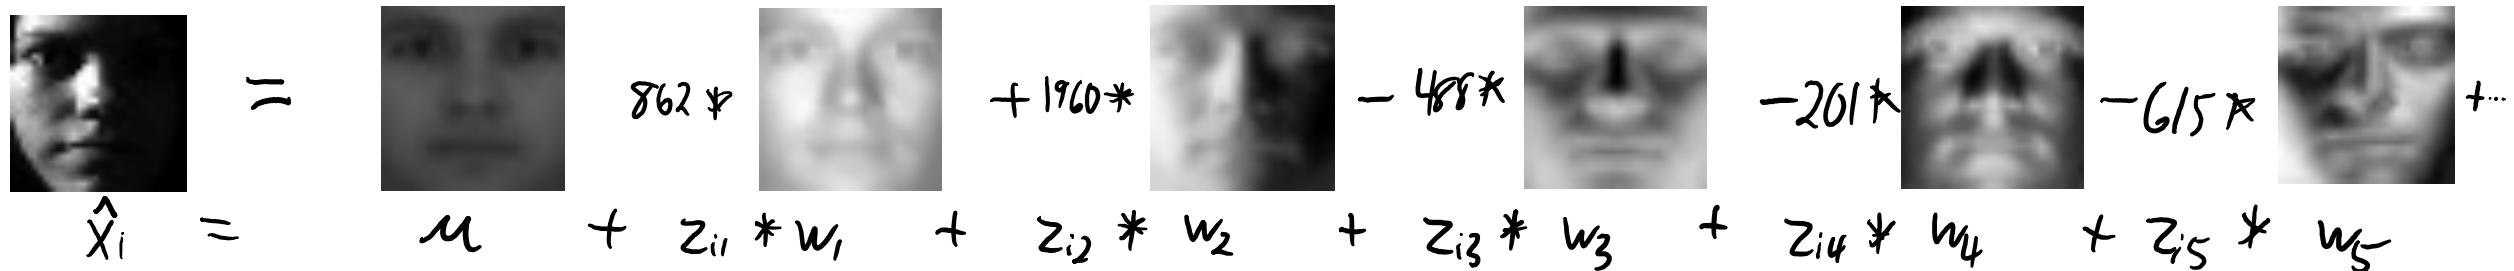
VQ vs. PCA vs. NMF

- How *should* we represent faces?
 - Vector quantization (k-means).
 - Replace face by the **average face in a cluster**.
 - ‘Grandmother cell’: one neuron = one face.
 - **Can’t distinguish between people** in the same cluster (only ‘ k ’ possible faces).
 - Almost certainly not true: too few neurons.

$$\hat{x}_i = 0 * \text{[face 1]} + 0 * \text{[face 2]} + 0 * \text{[face 3]} + 1 * \text{[face 4]} + 0 * \text{[face 5]} + 0 * \text{[face 6]} + \dots$$
$$\hat{x}_i = z_{i1} * w_1 + z_{i2} * w_2 + z_{i3} * w_3 + z_{i4} * w_4 + z_{i5} * w_5 + z_{i6} * w_6$$

VQ vs. PCA vs. NMF

- How *should* we represent faces?
 - Vector quantization (k-means).
 - PCA (orthogonal basis).
 - Global average plus linear combination of “eigenfaces”.
 - “Distributed representation”.
 - Coded by pattern of group of neurons: can represent infinite number of faces by changing z_i .
 - But “eigenfaces” are not intuitive ingredients for faces.
 - PCA tends to use positive/negative cancelling bases.

$$\hat{x}_i = \mu + z_{i1} * w_1 + z_{i2} * w_2 + z_{i3} * w_3 + z_{i4} * w_4 + z_{i5} * w_5 + \dots$$


VQ vs. PCA vs. NMF

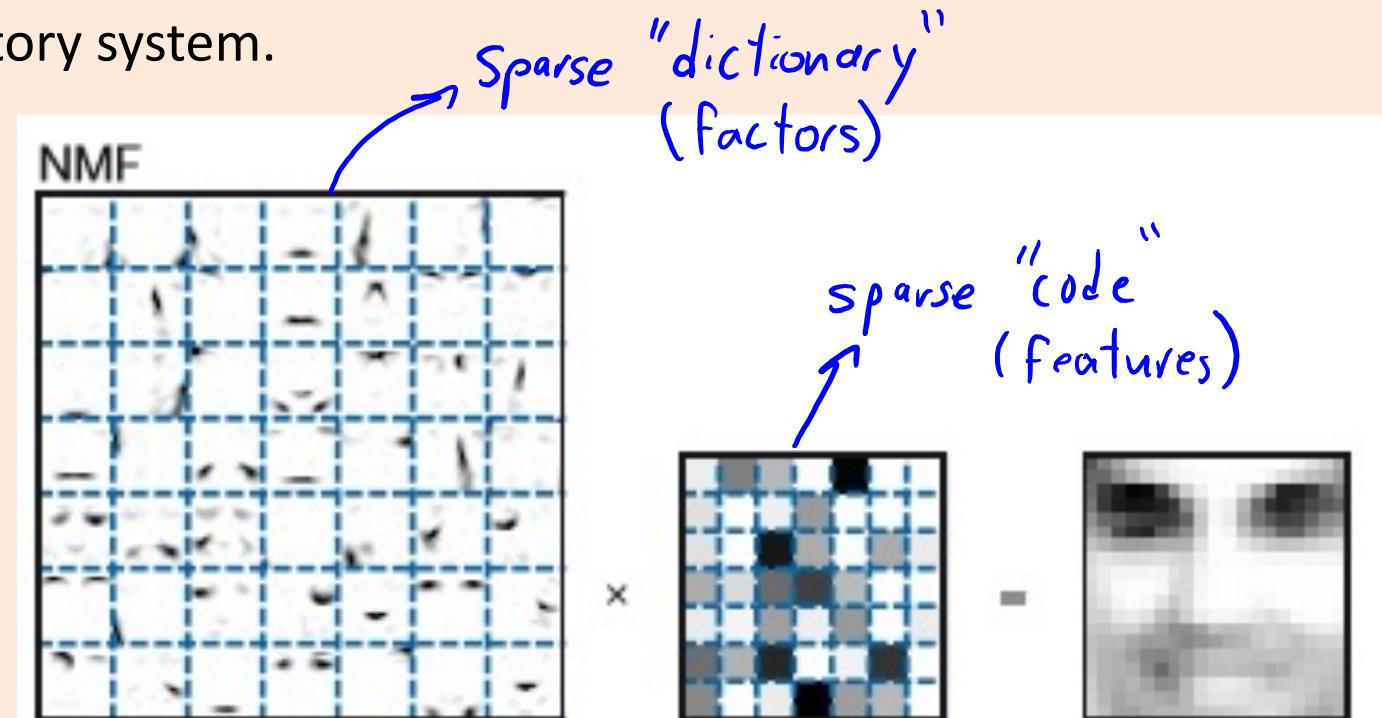
- How *should* we represent faces?
 - Vector quantization (k-means).
 - PCA (orthogonal basis).
 - NMF (non-negative matrix factorization):
 - Instead of orthogonality/ordering in W, require W and Z to be non-negativity.
 - Example of “sparse coding”:
 - The z_i are sparse so each face is coded by a small number of neurons.
 - The w_c are sparse so neurons tend to be “parts” of the object.

$$\begin{aligned} \text{Face Image} &= 4,2 * \text{Image}_1 + 0 * \text{Image}_2 + 0 * \text{Image}_3 + 3,3 * \text{Image}_4 + 0 * \text{Image}_5 + \dots \\ \hat{x}_i &= z_{i1} * w_1 + z_{i2} * w_2 + z_{i3} * w_3 + z_{i4} * w_4 + z_{i5} * w_5 \end{aligned}$$

bonus!

Representing Faces

- Why sparse coding?
 - “Parts” are intuitive, and brains seem to use sparse representation.
 - Energy efficiency if using sparse code.
 - Increase number of concepts you can memorize?
 - Some evidence in fruit fly olfactory system.



Warm-up to NMF: Non-Negative Least Squares

- Consider our usual **least squares** problem:

$$f(w) = \frac{1}{2} \sum_{i=1}^n (w^\top x_i - y_i)^2$$

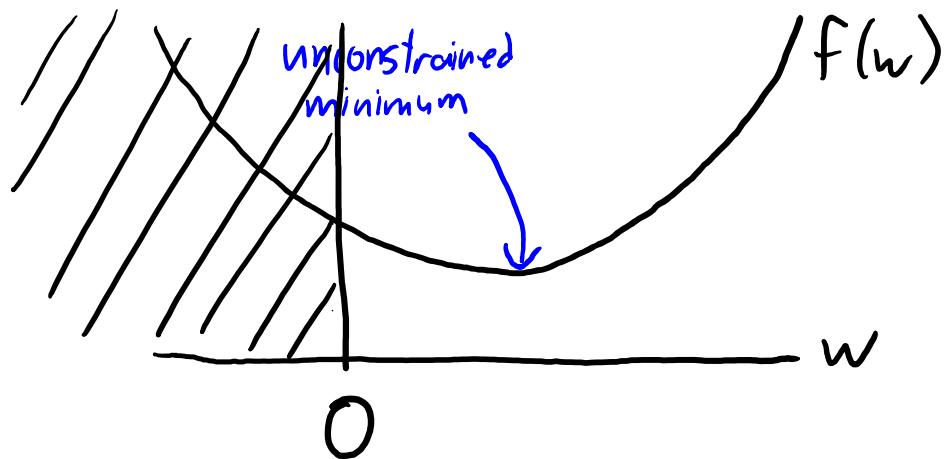
- But assume y_i and elements of x_i are **non-negative**:
 - Could be sizes ('height', 'milk', 'km') or counts ('vicodin', 'likes', 'retweets').
- Assume we want elements of ' w ' to be **non-negative**, too:
 - **No physical interpretation to negative weights.**
 - If x_{ij} is amount of product you produce, what does $w_j < 0$ mean?
- **Non-negativity leads to sparsity...**

Sparsity and Non-Negative Least Squares

- Consider 1D non-negative least squares objective:

$$f(w) = \frac{1}{2} \sum_{i=1}^n (w x_i - y_i)^2 \text{ with } w \geq 0$$

- Plotting the (constrained) objective function:



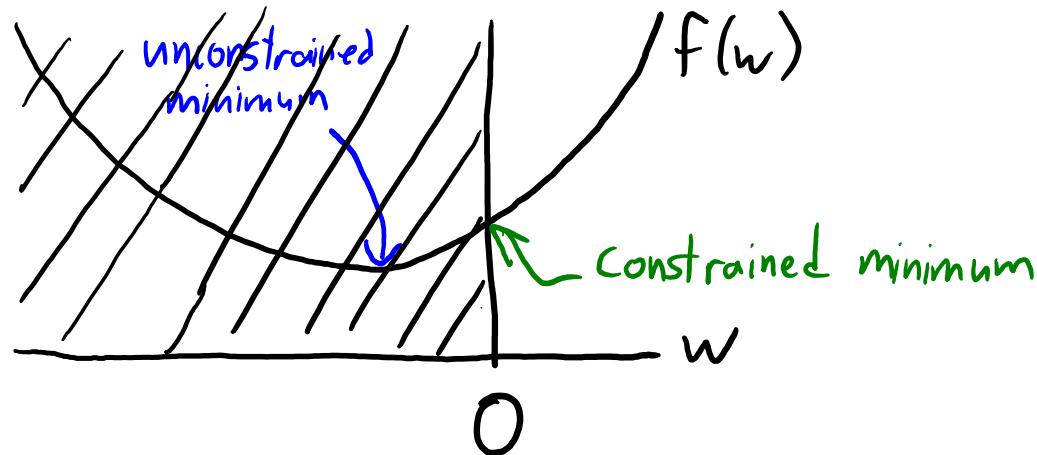
- In this case, non-negative solution is least squares solution.

Sparsity and Non-Negative Least Squares

- Consider 1D non-negative least squares objective:

$$f(w) = \frac{1}{2} \sum_{i=1}^n (w x_i - y_i)^2 \text{ with } w \geq 0$$

- Plotting the (constrained) objective function:



- In this case, non-negative solution is $w = 0$.

Sparsity and Non-Negativity

- Similar to L1-regularization, non-negativity leads to sparsity.
 - Also regularizes: w_j are smaller since can't "cancel" negative values.
 - Sparsity leads to cheaper predictions and often to more interpretability.
 - Non-negative weights are often also more interpretable.

- How can we minimize $f(w)$ with non-negative constraints?
 - Naive approach: solve least squares problem, set negative w_j to 0.

$$\text{Compute } w = (x^\top x)^{-1} (x^\top y)$$

$$\text{Set } w_j = \max\{0, w_j\}$$

- This is correct when $d = 1$.
- Can be worse than setting $w = 0$ when $d \geq 2$.

Sparsity and Non-Negativity

- Similar to L1-regularization, non-negativity leads to sparsity.
 - Also regularizes: w_j are smaller since can't "cancel" out negative values.
- How can we minimize $f(w)$ with non-negative constraints?
 - A correct approach is projected gradient algorithm:
 - Run a gradient descent iteration:

$$w^{t+1} = w^t - \alpha^t \nabla f(w^t)$$

- After each step, set negative values to 0.

$$w_j^{t+1} = \max\{0, w_j^{t+1}\}$$

- Repeat.

Sparsity and Non-Negativity

- Projected gradient algorithm:

$$w^{t+\frac{1}{2}} = w^t - \alpha^t \nabla f(w^t) \quad w_j^{t+1} = \max\{0, w_j^{t+\frac{1}{2}}\}$$

- Similar properties to gradient descent:

- Guaranteed decrease of ‘ f ’ if α^t is small enough.
- Reaches local minimum under weak assumptions (global minimum for convex ‘ f ’).
 - Least squares objective is still convex when restricted to non-negative variables.
- Solution is a “fixed point”: $w^* = \max\{0, w^* - \alpha^t \nabla f(w^*)\}$.
 - Use this to decide when to stop.

- A generalization is “proximal-gradient”:

- Instead of constraints, allows non-smooth terms (“findMinL1”).

Projected-Gradient for NMF

- Back to the **non-negative matrix factorization (NMF)** objective:

$$f(W, Z) = \sum_{i=1}^n \sum_{j=1}^d (w_i^j z_j - x_{ij})^2 \quad \text{with } w_{ij} \geq 0 \text{ and } z_j \geq 0$$

- Different ways to use **projected gradient**:

- Alternate between projected gradient steps on 'W' and on 'Z'.
- Or run projected gradient on both at once.
- Or sample a random 'i' and 'j' and do **stochastic projected gradient**.

Set $z_i^{t+1} = z_i^t - \alpha^t \nabla_{z_i} f(W, Z)$ and $(w_i^j)^{t+1} = (w_i^j)^t - \alpha^t \nabla_{w_i^j} f(W, Z)$ for selected i and j

- Non-convex** and (unlike PCA) is **sensitive to initialization**.

- Hard to find the global optimum.
- Typically use **random initialization**.
- Also, we **usually don't center the data** with NMF.

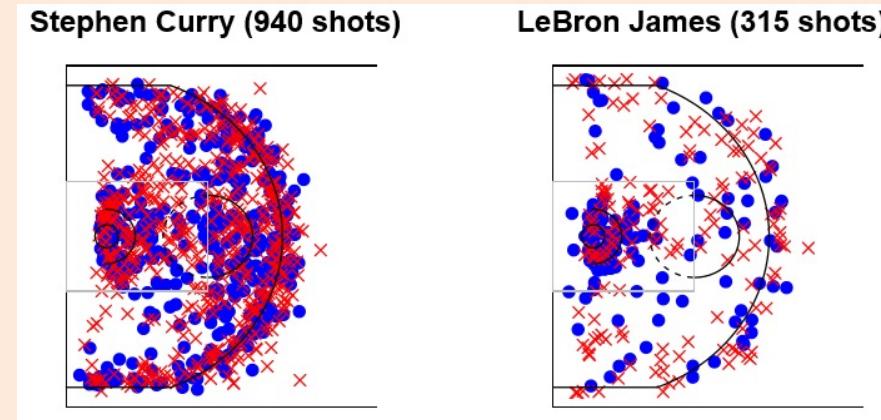
(Keep other values of W and Z fixed)

Then set negative values to 0.

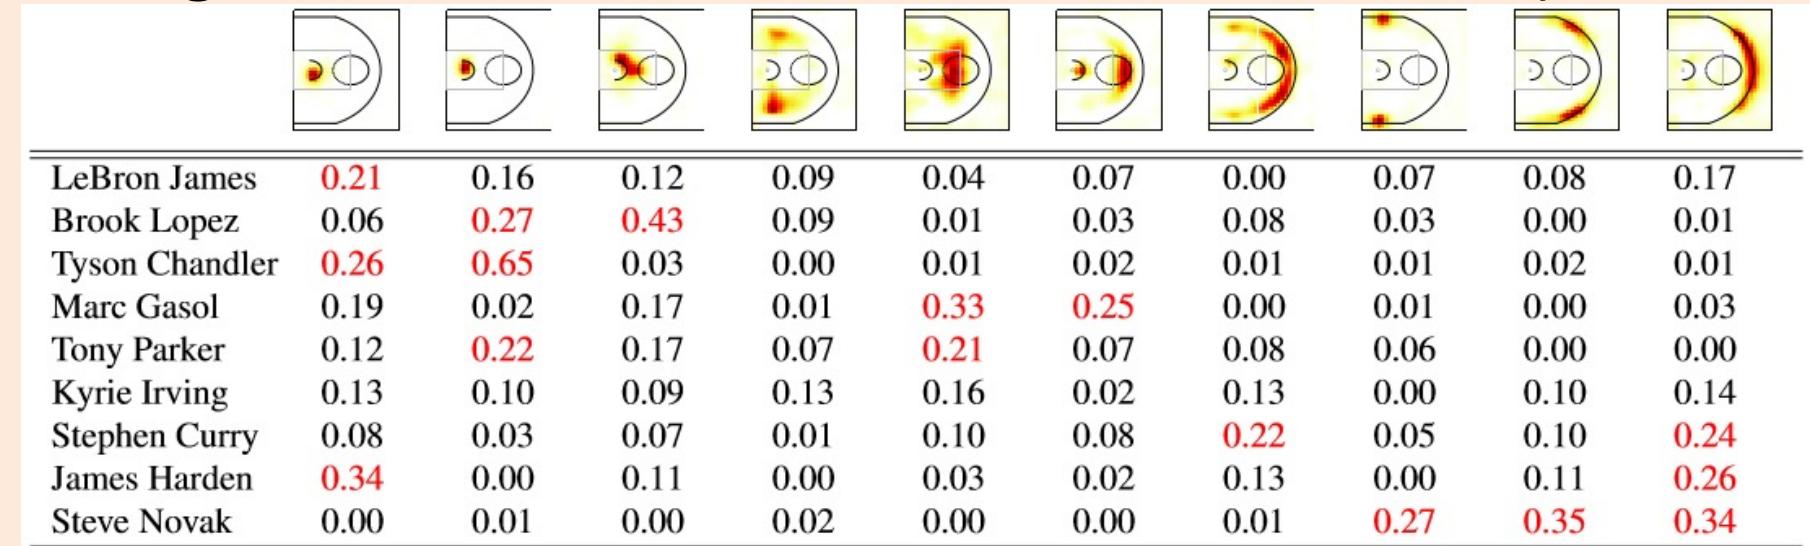
bonus!

Application: Sports Analytics

- NBA shot charts:



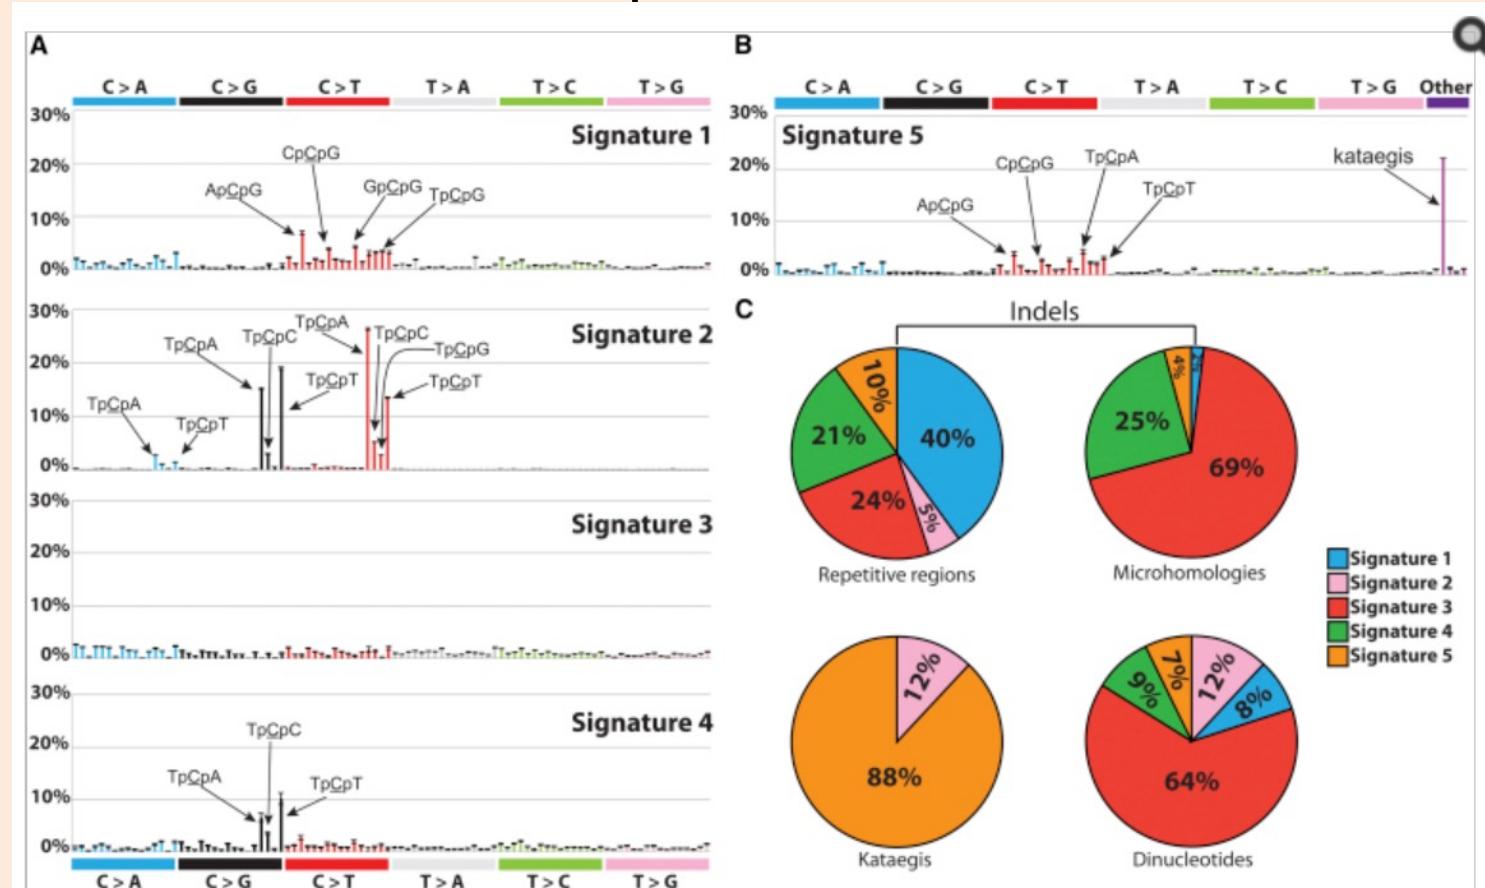
- NMF (using “KL divergence” loss with k=10 and smoothed data).
 - Negative values would not make sense here.



bonus!

Application: Cancer “Signatures”

- What are common sets of mutations in different cancers?
 - May lead to new treatment options.



(pause)

Beyond Squared Error

- Our objective for latent-factor models (LFM):

$$f(w, z) = \sum_{i=1}^n \sum_{j=1}^d (\langle w^j, z_i \rangle - x_{ij})^2$$

- As before, there are alternatives to squared error.

$$f(w, z) = \sum_{i=1}^n \sum_{j=1}^d \text{loss}(\langle w^j, z_i \rangle, x_{ij})$$

Error for predicting $\langle w^j, z_i \rangle$
when true value is x_{ij}

- If X consists of +1 and -1 values, we could use logistic loss:

$$f(w, z) = \sum_{i=1}^n \sum_{j=1}^d \log(1 + \exp(-x_{ij} \langle w^j, z_i \rangle))$$

Robust PCA

- Robust PCA methods use the absolute error:

$$f(w, z) = \sum_{i=1}^n \sum_{j=1}^d |\langle w_j^i z_i \rangle - x_{ij}|$$

- Will be robust to outliers in the matrix 'X'.
- Encourages "residuals" r_{ij} to be exactly zero.
– Non-zero r_{ij} are where the "outliers" are.

Applying robust PCA
to video frames



Regularized Matrix Factorization

- Recently people have also considered L2-regularized PCA:

$$f(W, Z) = \frac{1}{2} \|ZW - X\|_F^2 + \frac{\lambda_1}{2} \|W\|_F^2 + \frac{\lambda_2}{2} \|Z\|_F^2$$

- Replaces normalization/orthogonality/sequential-fitting.
 - Often gives lower reconstruction error on test data.
 - But requires regularization parameters λ_1 and λ_2 .
- Need to regularize W and Z because of scaling problem.
 - If you only regularize 'W' it doesn't do anything.
 - I could take unregularized solution, replace W by αW for a tiny α to shrink $\|W\|_F$ as much as I want, then multiply Z by $(1/\alpha)$ to get same solution.
 - Similarly, if you only regularize 'Z' it doesn't do anything.

Sparse Matrix Factorization

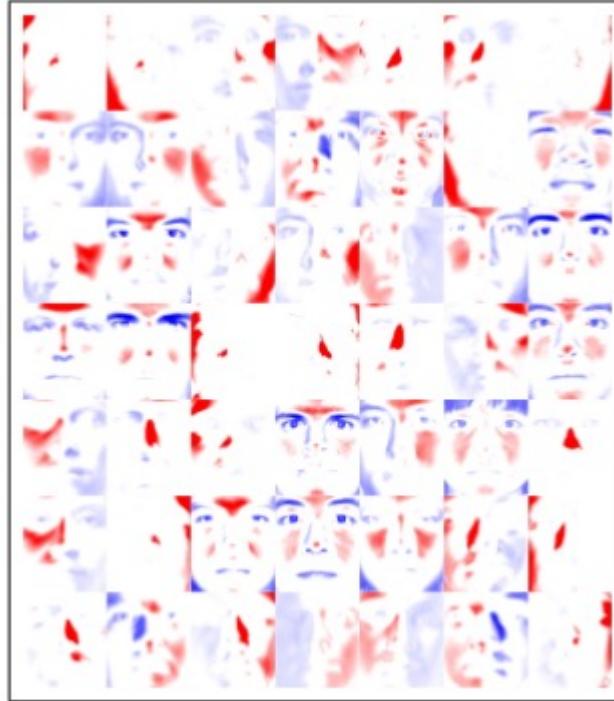
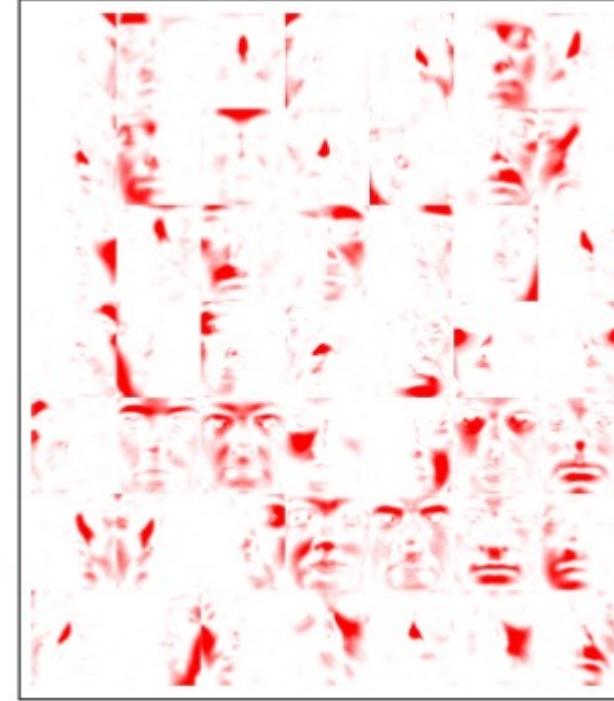
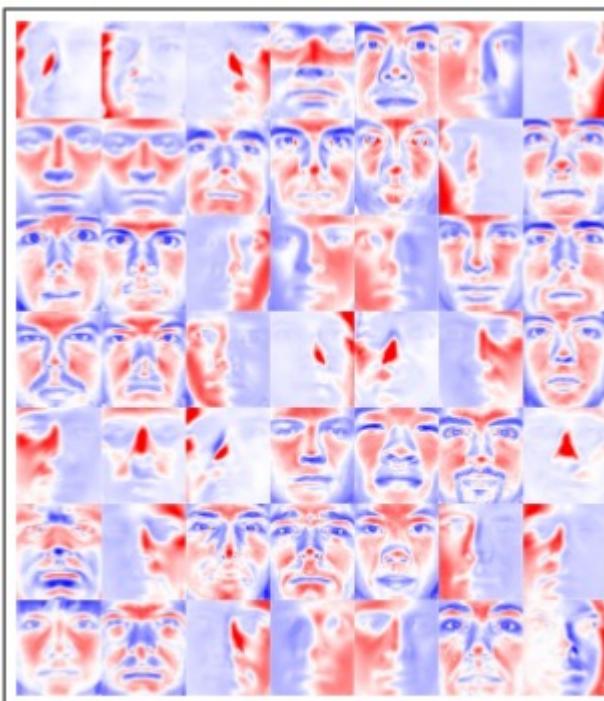
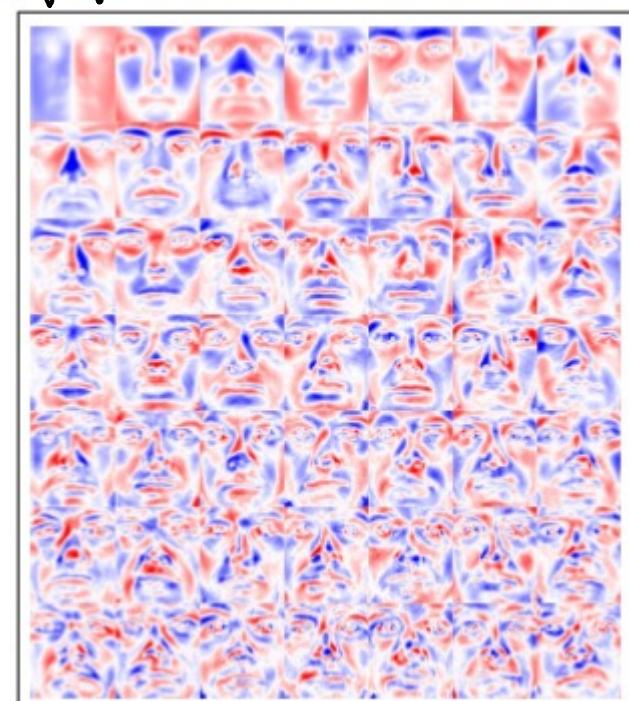
- Instead of non-negativity, we could use L1-regularization:

$$f(W, Z) = \frac{1}{2} \|ZW - X\|_F^2 + \frac{\lambda_1}{2} \sum_{i=1}^n \|z_i\|_1 + \frac{\lambda_2}{2} \sum_{j=1}^d \|w_j\|_1$$

- Called **sparse coding** (L1 on ‘Z’) or **sparse dictionary learning** (L1 on ‘W’).
- **Disadvantage of using L1-regularization** over non-negativity:
 - Sparsity controlled by λ_1 and λ_2 so you need to set these.
- **Advantage of using L1-regularization:**
 - Sparsity controlled by λ_1 and λ_2 , so you can **control amount of sparsity**.
 - Negative coefficients often do make sense.

Matrix Factorization with L1-Regularization

blue: negative
red: positive



PCA without orthogonality

sparsity due to
non-negativity

sparsity due to
L₁-regularization

bonus!

Sparse Matrix Factorization

- Instead of non-negativity, we could use L1-regularization:

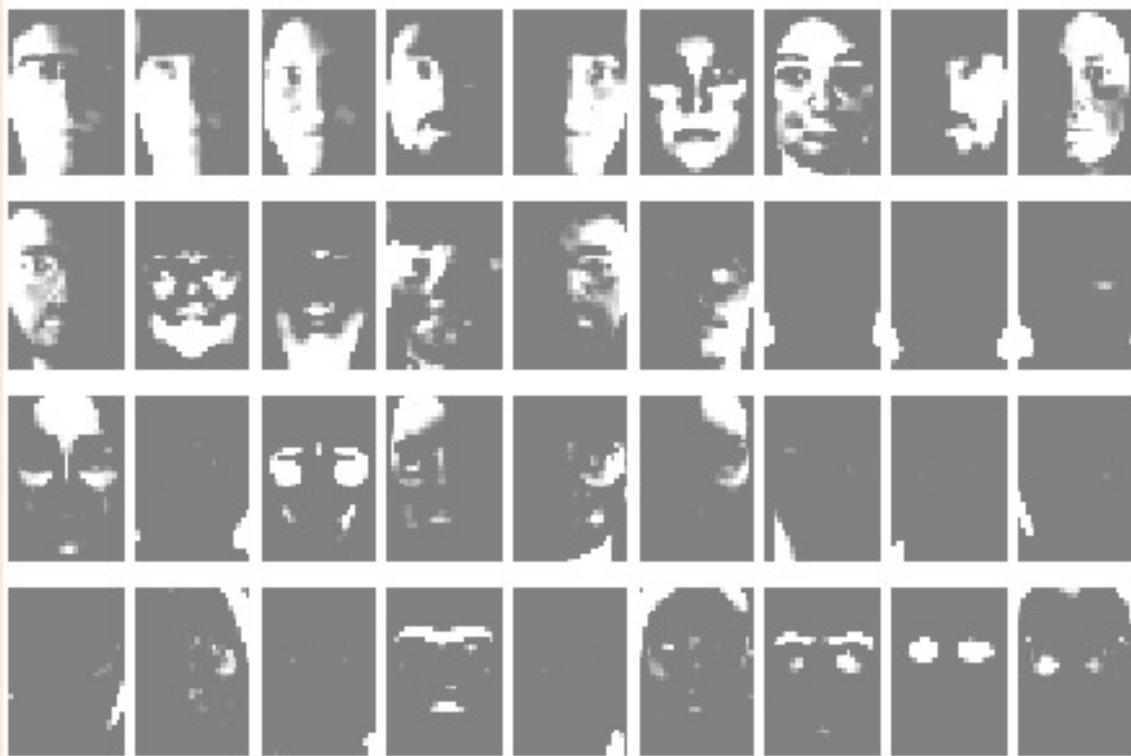
$$f(W, Z) = \frac{1}{2} \|ZW - X\|_F^2 + \frac{\lambda_1}{2} \sum_{i=1}^n \|z_i\|_1 + \frac{\lambda_2}{2} \sum_{j=1}^d \|w_j\|_1$$

- Called **sparse coding** (L1 on ‘Z’) or **sparse dictionary learning** (L1 on ‘W’).
- Many variations exist:
 - Mixing L2-regularization and L1-regularization.
 - Or normalizing ‘W’ (in L2-norm or L1-norm) and regularizing ‘Z’.
 - **K-SVD** constrains each z_i to have at most ‘k’ non-zeroes:
 - K-means is special case where $k = 1$.
 - PCA is special case where $k = d$.

bonus!

Recent Work: Structured Sparsity

- “Structured sparsity” considers dependencies in sparsity patterns.
 - Can enforce that “parts” are convex regions.



NMF



Sparse PCA with “structured” sparsity

Summary

- Biological motivation for orthogonal and/or sparse latent factors.
- Alternating minimization and stochastic gradient:
 - Iterative algorithms for minimizing PCA objective.
- Non-negative matrix factorization: LFM with no negative values.
 - Non-negativity constraints lead to sparse solution.
 - Projected gradient adds constraints to gradient descent.
- Many of our regression tricks can be used with LFMs:
 - Robust PCA uses absolute error to be robust to outliers.
 - L1-regularization leads to sparse factors/weights.
- Next time: the million-dollar NetFlix challenge.

Proof: “Synthesis” View = “Analysis” View ($WW^T = I$)

bonus!

- The variance of the z_{ij} (maximized in “analysis” view):

$$\begin{aligned}
 \frac{1}{n^k} \sum_{i=1}^n \|z_i - \mu_z\|^2 &= \frac{1}{n^k} \sum_{i=1}^n \|Wx_i\|^2 \quad (\mu_z = 0 \text{ and } z_i = Wx_i \text{ if } \|W_i\| = 1 \text{ and } W_i^T W_i = 0) \\
 &= \frac{1}{n^k} \sum_{i=1}^n x_i^T W^T W x_i = \frac{1}{n^k} \sum_{i=1}^n \text{Tr}(x_i^T W^T W x_i) = \frac{1}{n^k} \sum_{i=1}^n \text{Tr}(W^T W x_i x_i^T) \\
 &= \frac{1}{n^k} \text{Tr}(W^T W \sum_{i=1}^n x_i x_i^T) = \frac{1}{n^k} \text{Tr}(W^T W X^T X)
 \end{aligned}$$

linearity of trace “cyclic” property of trace

- The distance to the hyper-plane (minimized in “synthesis” view):

$$\begin{aligned}
 \|2W - X\|_F^2 &= \|XW^T W - X\|_F^2 = \text{Tr}((XW^T W - X)^T (XW^T W - X)) \\
 &= \text{Tr}(W^T W X^T X W^T W) - 2\text{Tr}(W^T W X^T X) + \text{Tr}(X^T X) \\
 &= \text{Tr}(W^T \underbrace{W W^T}_I W X^T X) - 2\text{Tr}(W^T W X^T X) + \text{Tr}(X^T X) \\
 &= -\text{Tr}(W^T W X^T X) + (\text{constant})
 \end{aligned}$$

$\|A\|_F^2 = \text{Tr}(A^T A)$ Solved by same ‘W’
 $= XW^T$ ↑

bonus!

Canonical Correlation Analysis (CCA)

- Suppose we have two matrices, ‘X’ and ‘Y’.
- Want to find matrices W_X and W_Y that maximize correlation.
 - “What are the latent factors in common between these datasets?”
- Define the correlation matrices:

$$\Sigma_{xx} = \frac{1}{n} \sum_{i=1}^n x_i x_i^\top \quad \Sigma_{yy} = \frac{1}{n} \sum_{i=1}^n y_i y_i^\top \quad \Sigma_{xy} = \frac{1}{n} \sum_{i=1}^n x_i y_i^\top$$

- Canonical correlation analysis (CCA) maximizes

$$\text{Tr}(W_Y^\top W_X \Sigma_{xx}^{-\frac{1}{2}} \Sigma_{xy} \Sigma_{yy}^{-\frac{1}{2}})$$

- Subject to W_X and W_Y having orthogonal rows.
- Computationally, equivalent to PCA with a different matrix.
 - Using the “analysis” view that PCA maximizes $\text{Tr}(W^\top W X^\top X)$.

bonus!

Kernel PCA

- From the “analysis” view (with orthogonal PCs) PCA maximizes:

$$\text{Tr}(W^T W X^T X)$$

- It can be shown that the solution has the form (see [here](#)):

$$W = UX$$

$\underbrace{W}_{k \times d} \quad \underbrace{U}_{k \times n} \underbrace{X}_{n \times 1}$

- Re-parameterizing in terms of ‘U’ gives a [kernelized PCA](#):

$$\text{Tr}(X^T V^T V X X^T X) = \text{Tr}(V^T V X X^T X X^T)$$

$\underbrace{V}_{K \times k} \quad \underbrace{X}_{k \times k} \quad \underbrace{X^T}_{k \times k}$

- It’s hard to initially center data in ‘Z’ space, but you can [form the centered kernel matrix](#) (see [here](#)).

bonus!

Probabilistic PCA

- With zero-mean (“centered”) data, in PCA we assume that

$$x_i \approx W^T z_i$$

- In **probabilistic PCA** we assume that

$$x_i \sim N(W^T z_i, \sigma^2 I) \quad z_i \sim N(0, I)$$

- Integrating over ‘Z’ the marginal likelihood given ‘W’ is Gaussian,

$$x_i | W \sim N(0, W^T W + \sigma^2 I)$$

- Regular PCA is obtained as the limit of σ^2 going to 0.

bonus!

Generalizations of Probabilistic PCA

- Probabilistic PCA model:

$$x_i \mid W \sim \mathcal{N}(0, W^T W + \sigma^2 I)$$

- Why do we need a probabilistic interpretation?
- Shows that PCA fits a Gaussian with restricted covariance.
 - Hope is that $W^T W + \sigma^2 I$ is a good approximation of $X^T X$.
- Gives precise connection between PCA and factor analysis.

bonus!

Factor Analysis

- Factor analysis is a method for discovering latent factors.
- Historical applications are measures of intelligence and personality.

Trait	Description
O penness	Being curious, original, intellectual, creative, and open to new ideas.
C onscientiousness	Being organized, systematic, punctual, achievement-oriented, and dependable.
E xtraversion	Being outgoing, talkative, sociable, and enjoying social situations.
A greeableness	Being affable, tolerant, sensitive, trusting, kind, and warm.
N euroticism	Being anxious, irritable, temperamental, and moody.

- A standard tool and widely-used across science and engineering.

bonus!

PCA vs. Factor Analysis

- PCA and FA both write the matrix ‘X’ as

$$X \approx ZW$$

- PCA and FA are both based on a Gaussian assumption.
- Are PCA and FA the same?
 - Both are more than 100 years old.
 - People are still arguing about whether they are the same:
 - Doesn't help that some packages run PCA when you call their FA method.

bonus!

Google pca vs. factor analysis

All Images Videos News Maps More Search tools

About 358,000 results (0.17 seconds)

[PDF] Principal Component Analysis versus Exploratory Factor ...
www2.sas.com/proceedings/sugi30/203-30.pdf ▾
by DD Suhr - Cited by 118 - Related articles
1. Paper 203-30. Principal Component Analysis vs. Exploratory Factor Analysis.
Diana D. Suhr, Ph.D. University of Northern Colorado. Abstract. Principal ...

pca - What are the differences between Factor Analysis and ...
stats.stackexchange.com/.../what-are-the-differences-between-factor-anal... ▾
Aug 12, 2010 - Principal Component Analysis (PCA) and Common Factor Analysis (CFA) differently one has to interpret the strength of loadings in PCA vs.

What are the differences between principal components ...
support.minitab.com/.../factor-analysis/differences-between-pca-and-facto... ▾
Principal Components Analysis and Factor Analysis are similar because both procedures are used to simplify the structure of a set of variables. However, the ...

[PDF] Principal Components Analysis - UNT
<https://www.unt.edu/rss/class/.../Principal%20Components%20Analysis.p...> ▾
PCA vs. Factor Analysis. • It is easy to make the mistake in assuming that these are the same techniques, though in some ways exploratory factor analysis and ...

Factor analysis versus Principal Components Analysis (PCA)
psych.wisc.edu/henriques/pca.html ▾
Jun 19, 2010 - Factor analysis versus PCA. These techniques are typically used to analyze groups of correlated variables representing one or more common ...

[PDF] Principal Component Analysis and Factor Analysis
www.stats.ox.ac.uk/~ripley/MultAnal_HT2007/PC-FA.pdf ▾
where D is diagonal with non-negative and decreasing values and U and V Factor analysis and PCA are often confused, and indeed SPSS has PCA as.

How can I decide between using principal components ...
https://www.researchgate.net/.../How_can_I_decide_between_using_prin... ▾
Factor analysis (FA) is a group of statistical methods used to understand and simplify patterns ... Retrieved from <http://pareonline.net/getvn.asp?v=10&n=7> ... Principal component analysis (PCA) is a method of factor extraction (the second step ...

[PDF] Exploratory Factor Analysis and Principal Component An...
www.lesahoffman.com/948/948_Lecture2_EFA_PCA.pdf ▾
2 very different schools of thought on exploratory factor analysis (EFA) vs. principal components analysis (PCA): ➤ EFA and PCA are TWO ENTIRELY ...

Factor analysis - Wikipedia, the free encyclopedia
https://en.wikipedia.org/wiki/Factor_analysis ▾
Jump to Exploratory **factor analysis** versus principal components ... - [edit]. See also: Principal component analysis and Exploratory factor analysis.

[PDF] The Truth about PCA and Factor Analysis
www.stat.cmu.edu/~cshalizi/350/lectures/13/lecture-13.pdf ▾
Sep 28, 2009 - nents and factor analysis, we'll wrap up by looking at their uses and

PCA vs. Factor Analysis

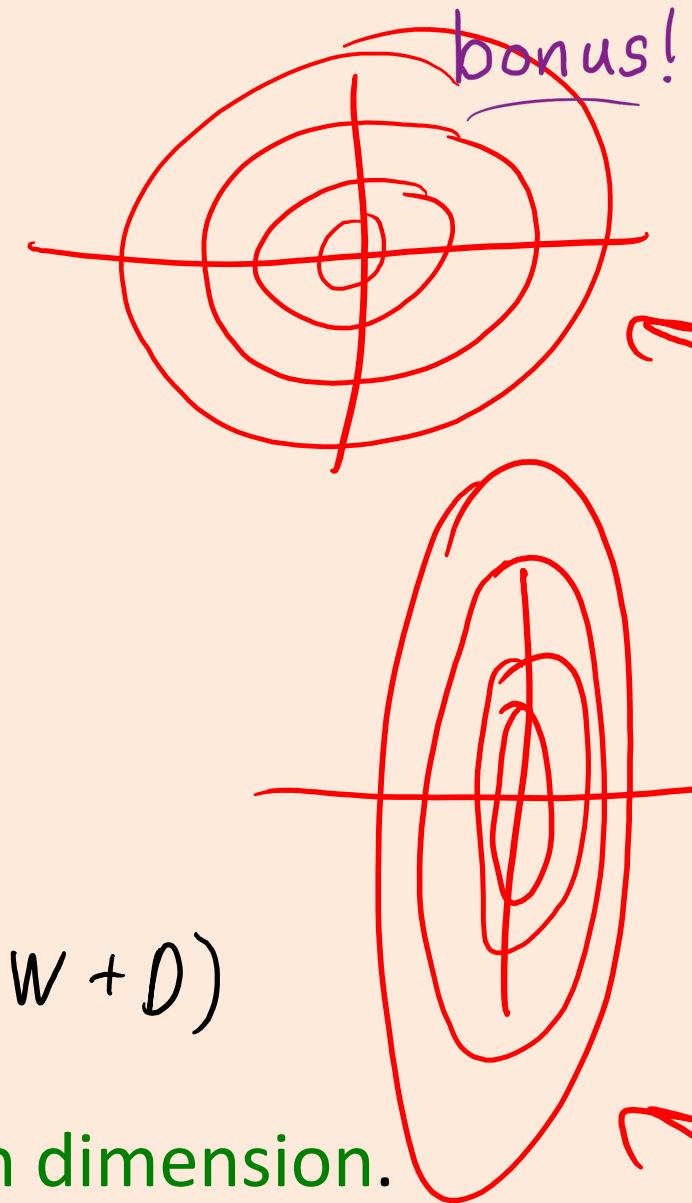
- In probabilistic PCA we assume:

$$x_i \sim \mathcal{N}(W^\top z_i, \sigma^2 I)$$

- In FA we assume for a diagonal matrix D that:

$$x_i \sim \mathcal{N}(W^\top z_i, D)$$

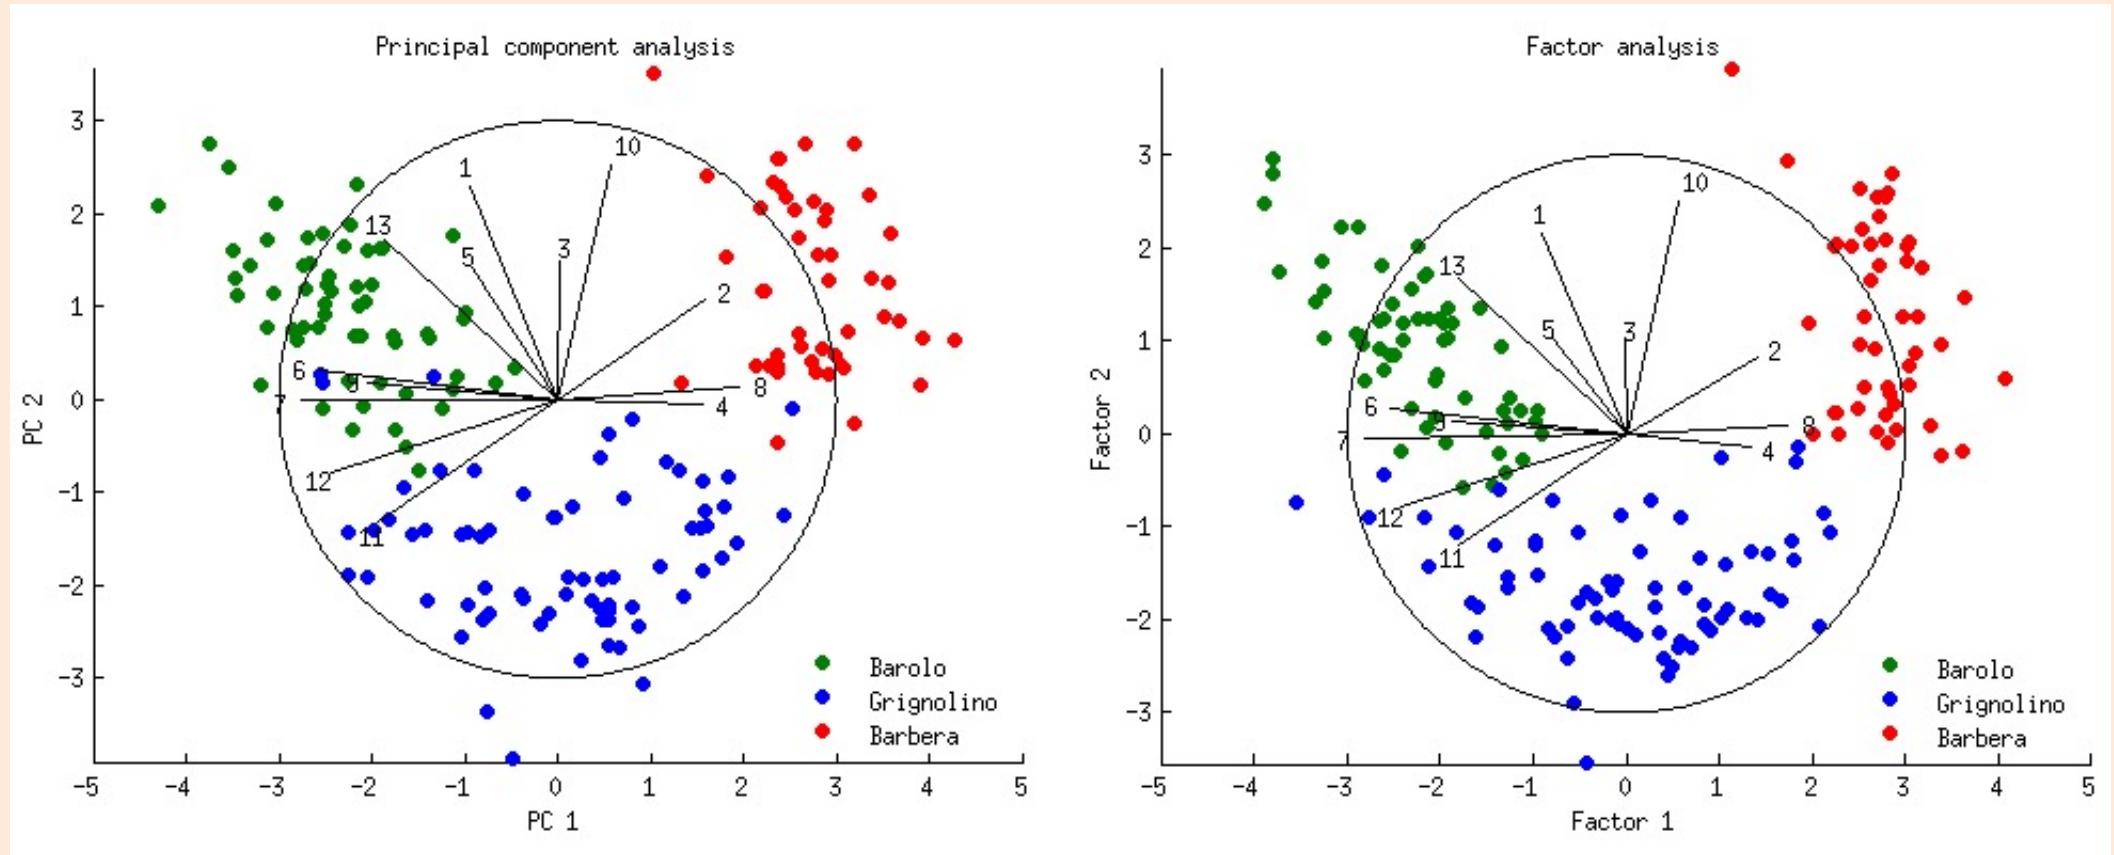
- The posterior in this case is: $x_i | W \sim \mathcal{N}(0, W^\top W + D)$
- The difference is you have a noise variance for each dimension.
 - FA has extra degrees of freedom.



bonus!

PCA vs. Factor Analysis

- In practice there often isn't a huge difference:



bonus!

Factor Analysis Discussion

- Differences with PCA:
 - Unlike PCA, FA is not affected by scaling individual features.
 - But unlike PCA, it's affected by rotation of the data.
 - No nice “SVD” approach for FA, you can get different local optima.
- Similar to PCA, FA is invariant to rotation of ‘W’.
 - So as with PCA you can't interpret multiple factors as being unique.

bonus!

Motivation for ICA

- Factor analysis has found an enormous number of applications.
 - People really want to find the “hidden factors” that make up their data.
- But PCA and FA **can't identify the factors.**

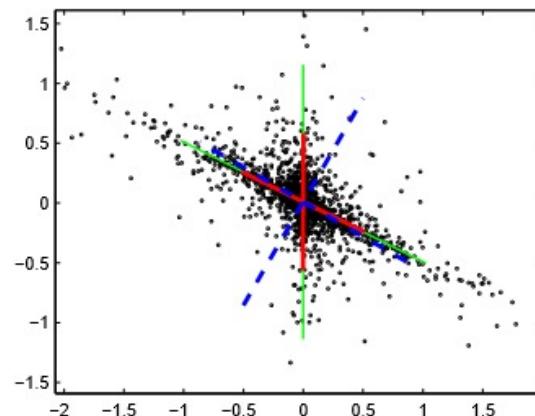


Figure : Latent data is sampled from the prior $p(x_i) \propto \exp(-5\sqrt{|x_i|})$ with the mixing matrix A shown in green to create the observed two dimensional vectors $y = Ax$. The red lines are the mixing matrix estimated by `ica.m` based on the observations. For comparison, PCA produces the blue (dashed) components. Note that the components have been scaled to improve visualisation. As expected, PCA finds the orthogonal directions of maximal variation. ICA however, correctly estimates the directions in which the components were independently generated.

bonus!

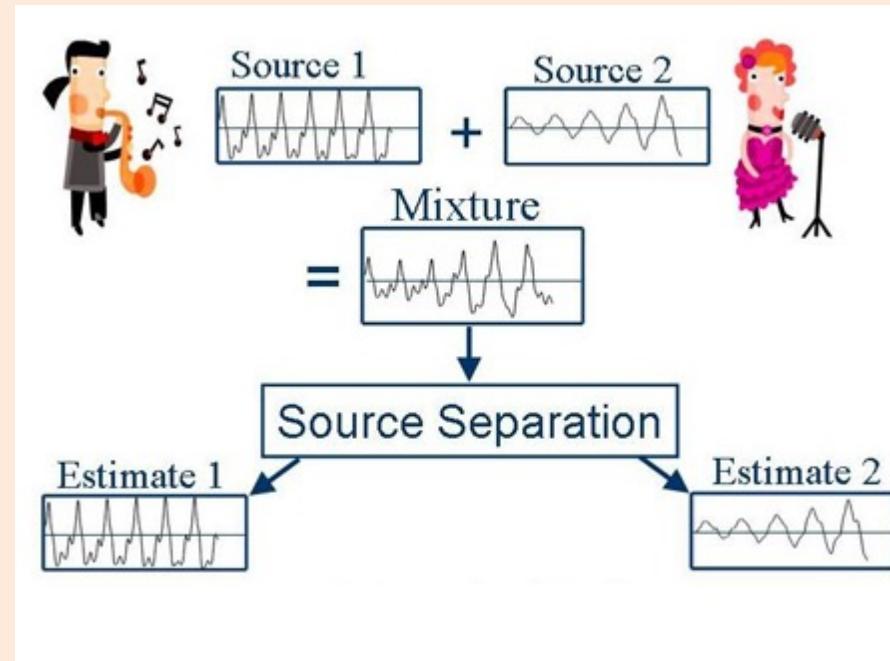
Motivation for ICA

- Factor analysis has found an enormous number of applications.
 - People really want to find the “hidden factors” that make up their data.
- But PCA and FA **can't identify the factors**.
 - We can rotate W and obtain the same model.
- **Independent component analysis (ICA)** is a more recent approach.
 - Around 30 years old instead of > 100.
 - Under certain assumptions it can **identify factors**.
- The canonical application of ICA is **blind source separation**.

bonus!

Blind Source Separation

- Input to blind source separation:
 - Multiple microphones recording multiple sources.



- Each microphone gets different mixture of the sources.
 - Goal is reconstruct sources (factors) from the measurements.

Independent Component Analysis Applications

bonus!

- ICA is replacing PCA and FA in many applications:

Some ICA applications are listed below:^[1]

- optical Imaging of neurons^[17]
- neuronal spike sorting^[18]
- face recognition^[19]
- modeling receptive fields of primary visual neurons^[20]
- predicting stock market prices^[21]
- mobile phone communications^[22]
- color based detection of the ripeness of tomatoes^[23]
- removing artifacts, such as eye blinks, from EEG data.^[24]

- Recent work shows that ICA can often resolve direction of causality.

bonus!

Limitations of Matrix Factorization

- ICA is a **matrix factorization** method like PCA/FA,

$$X = ZW$$

- Let's assume that $X = ZW$ for a "true" W with $k = d$.
 - Different from PCA where we assume $k \leq d$.
- There are only **3 issues** stopping us from finding "true" W .

bonus!

3 Sources of Matrix Factorization Non-Uniqueness

- **Label switching:** get same model if we **permute rows** of W.
 - We can exchange row 1 and 2 of W (and same columns of Z).
 - Not a problem because we don't care about order of factors.
- **Scaling:** get same model if you **scale a row**.
 - If we multiply row 1 of W by α , could multiply column 1 of Z by $1/\alpha$.
 - Can't identify sign/scale, but might hope to identify direction.
- **Rotation:** get same model if we **rotate W**.
 - Rotations correspond to orthogonal matrices Q, such matrices have $Q^TQ = I$.
 - If we rotate W with Q, then we have $(QW)^TQW = W^TQ^TQW = W^TW$.
- **If we could address rotation, we could identify the “true” directions.**

bonus!

A Unique Gaussian Property

- Consider an **independent prior** on each latent features z_c .
 - E.g., in PPCA and FA we use $N(0,1)$ for each z_c .
- If prior $p(z)$ is independent and **rotation-invariant** ($p(Qz) = p(z)$), then it must be Gaussian (only Gaussians have this property).
- The (non-intuitive) magic behind ICA:
 - If the priors are all **non-Gaussian**, it **isn't rotationally symmetric**.
 - In this case, we can **identify factors W** (up to permutations and scalings).

bonus!

PCA vs. ICA

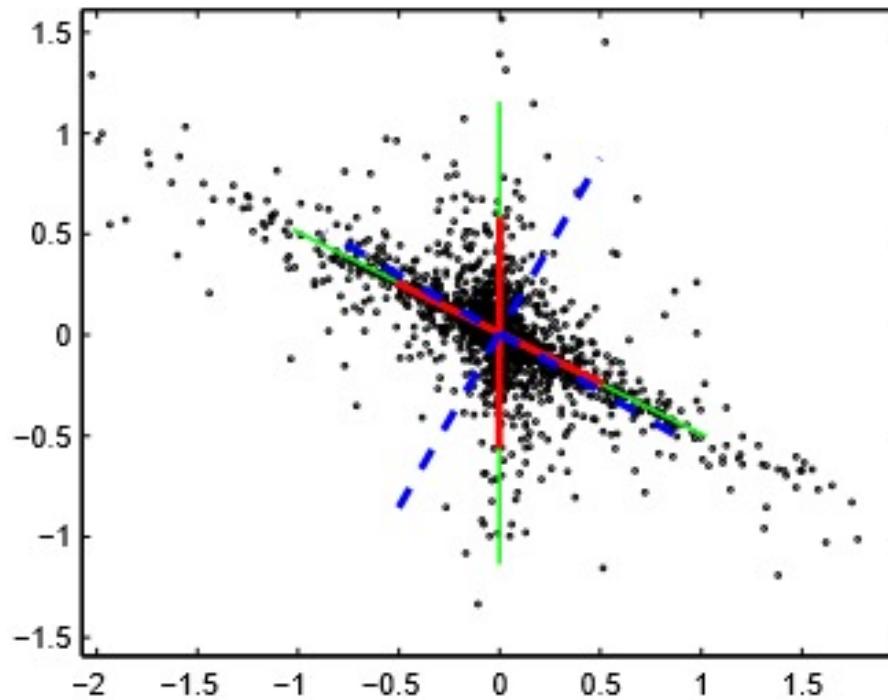


Figure : Latent data is sampled from the prior $p(x_i) \propto \exp(-5\sqrt{|x_i|})$ with the mixing matrix A shown in green to create the observed two dimensional vectors $y = Ax$. The red lines are the mixing matrix estimated by `ica.m` based on the observations. For comparison, PCA produces the blue (dashed) components. Note that the components have been scaled to improve visualisation. As expected, PCA finds the orthogonal directions of maximal variation. ICA however, correctly estimates the directions in which the components were independently generated.

bonus!

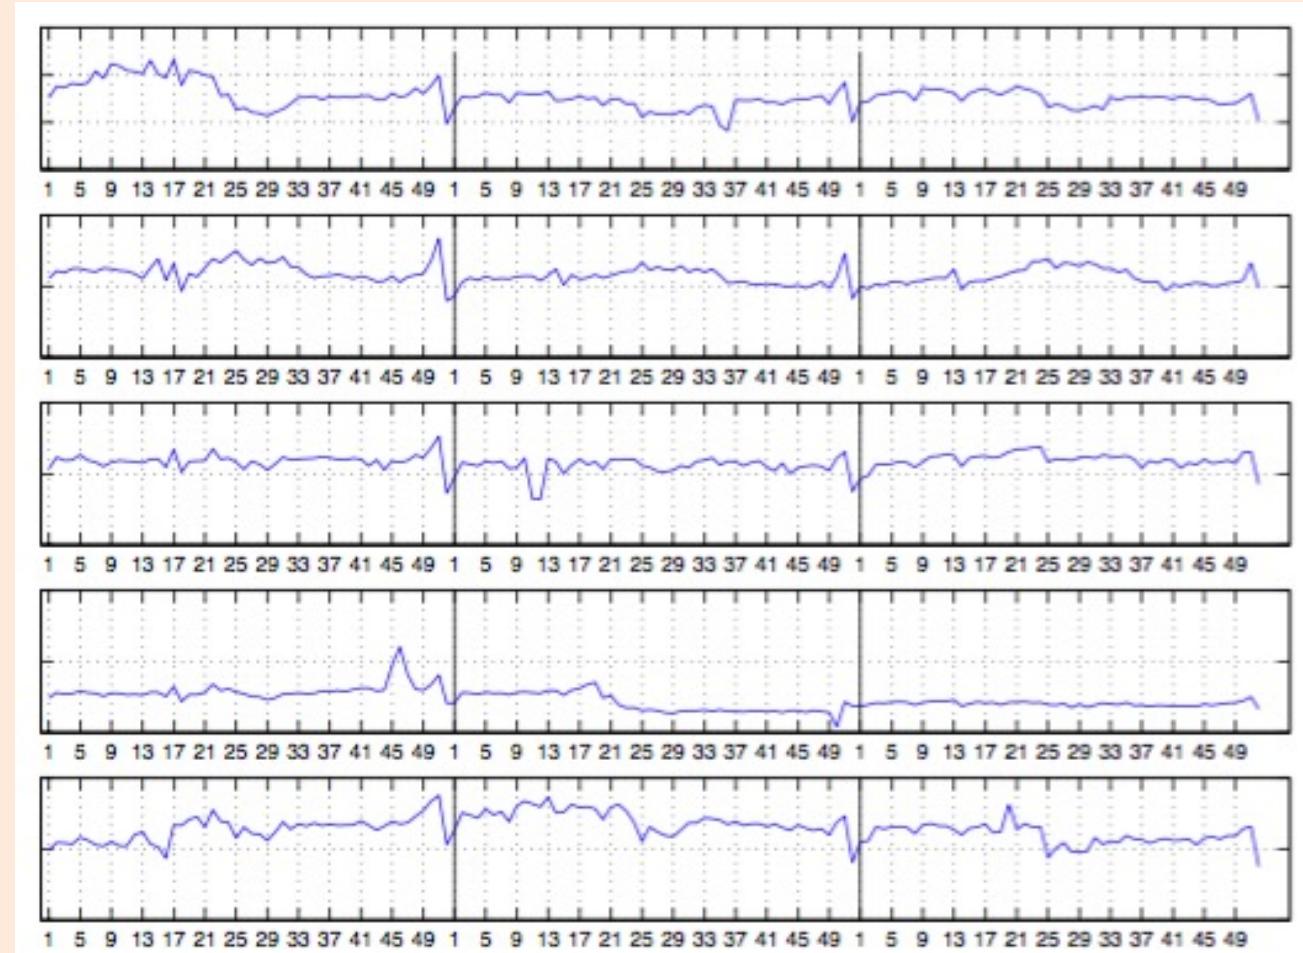
Independent Component Analysis

- In ICA we approximate X with ZW , assuming $p(z_{ic})$ are **non-Gaussian**.
- Usually we “center” and “whiten” the data before applying ICA.
- There are several penalties that encourage non-Gaussianity:
 - Penalize low **kurtosis**, since kurtosis is minimized by Gaussians.
 - Penalize high **entropy**, since entropy is maximized by Gaussians.
- The **fastICA** is a popular method maximizing kurtosis.

bonus!

ICA on Retail Purchase Data

- Cash flow from 5 stores over 3 years:



bonus!

ICA on Retail Purchase Data

- Factors found using ICA:

