


CPSC 340: Machine Learning and Data Mining

Convolutions


Fall 2021

Motivation: Automatic Brain Tumor Segmentation


- Task: labeling tumors and normal tissue in multi-modal MRI data.



Input:



Output:



- Applications:


- Radiation therapy target planning, quantifying treatment responses.
 - Mining growth patterns, image-guided surgery.

- Challenges:

- Variety of tumor appearances, similarity to normal tissue.
 - “You are never going to solve this problem.”

Naïve Voxel-Level Classifier

- We could treat classifying a voxel as **supervised learning**:
 - Standard representation of image: each pixel gets “intensity” between 0 and 255.




$$x_i = (98, 187, 246)$$

$y_i = \text{"tumour"}$

- We can formulate predicting y_i given x_i as supervised learning.
- But it **doesn't work** at all with these features.

Need to Summarize Local Context

- The individual pixel intensity values are almost meaningless:
 - The same x_i could lead to different y_i .



- Intensities not standardized.
- Non-trivial overlap in signal for different tissue types.
- “Partial volume” effects at boundaries of tissue types.


Need to Summarize Local Context

- We need to represent the “context” of the pixel (what is around it).



- Include all the values of neighbouring pixels as extra features?
 - Run into coupon collection problems: requires lots of data to find patterns.
- Measure neighbourhood summary statistics (mean, variance, histogram)?
 - Variation on bag of words problem: loses spatial information present in voxels.
- Standard approach uses convolutions to represent neighbourhood.

Example: Measuring “brightness” of an Area




- This pixel is in a “bright” area of the image, which reflects “bleeding” of tumour.
- But the actual numeric intensity value of the pixel is the **same as in darker “gray matter” areas.**
- I want a feature saying “this pixel is in a bright area of the image”.
 - This will help us identify that it’s a tumour pixel.

- How to measure brightness in area? Easy way: take **average pixel intensity** in “neighbourhood”.

$$z = \frac{1}{|nei|} \sum_{k \in nei} x_k$$


new feature is average value in neighbourhood.

- Applying this “averaging” to **every pixel** gives a new image:
- We can use “pixel value in new image” as a new feature.
 - New feature helps identify if pixel is in a “bright” area.




The annoying thing about squares

- “Take the average of a square window” loses spatial information.
- Example:



replace each pixel by the
average value of "window"

Take
average



Weird stuff

Pixels far away from “edge” become bright.


Noticeable features are “averaged out”

-Average is higher, but
location is lost.

“Park because line is
“surrounded” by dark areas

Fixing the “square” issues

- Consider instead “blurring” the image.
 - Gets rid of “local” noise, but better preserves spatial information.



- How do you “blur”?
 - Take **weighted average** of window, putting more “weight” on “close” pixels:

$$z = \sum w_k x_k$$


$k \in \text{ nei}$ \hookrightarrow weight on pixel x_k (averaging is special case where all pixels get equal weight)

Fixing the “square” issues

- Another neat thing we can do: use negative weights.
 - These features can describe “differences” across space.



“Average”
with
positive
and
negative
weights




Useful feature: “My neighbours are brighter than me”

- Taking a “weighted average of neighbours” is called “convolution”.
 - Gives you something like the “words” that make up image regions.

Convolution: Big Picture

- How do you use convolution to get features?
 - Apply **several different convolutions** to your image.
 - Each convolution gives a different “image” value at each location.
 - Use **theses different image values** to give features at each location.



Convolutions: Big Picture



- What can features coming from convolutions represent?
 - Some filters give you an **average value of the neighbourhood**.
 - Some filters **approximate the “first derivative”** in the neighbourhood.
 - “Is there a change from low to dark to bright?”
 - “If so, from which direction in space?”
 - Some filters **approximate the “second derivative”** in the neighbourhood.
 - “Is there a spike or is the change speeding up?”
- Hope: we can characterize **“what happens in a neighbourhood”**, with just a few numbers.

1D Convolution Example

- Consider a 1D “signal” (maybe from sound):
 - We’ll come back to images later.
- For each “time”:
 - Compute dot-product of signal at surrounding times with a “filter” of weights.

$$w = [-0.1416 \ -0.1781 \ -0.2746 \ 0.1640 \ 0.8607 \ 0.1640 \ -0.2746 \ -0.1781 \ -0.1416]$$

- This gives a new “signal”:
 - Measures a property of “neighbourhood”.
 - This particular filter shows a local “how spiky” value.



1D Convolution (notation is specific to this lecture)

- **1D convolution** input:

- Signal ‘x’ which is a vector length ‘n’.

- Indexed by $i = 1, 2, \dots, n$

- Filter ‘w’ which is a vector of length ‘ $2m+1$ ’:

- Indexed by $i = -m, -m+1, \dots, -2, 0, 1, 2, \dots, m-1, m$

$$x = [0 \ 1 \ 1 \ 2 \ 3 \ 5 \ 8 \ 13]$$

$$w = [0 \ -1 \ 2 \ -1 \ 0]$$

$w_{-2} \quad w_{-1} \quad w_0 \quad w_1 \quad w_2$

- Output is a vector of length ‘n’ with elements:

$$z_i = \sum_{j=-m}^m w_j x_{i+j}$$

- You can think of this as **centering w at position ‘i’**,
and **taking a dot product of ‘w’ with that “part” x_i .**

1D Convolution

- 1D convolution example:

- Signal ‘x’:

0	1	1	2	3	5	8	13
---	---	---	---	---	---	---	----

- Filter ‘w’:

0	-1	2	-1	0
---	----	---	----	---

- Convolution ‘z’:

--	--	--	--	--	--	--	--

1D Convolution

- 1D convolution example:

- Signal 'x':



- Filter 'w':



- Convolution 'z':



take dot-product
$$(0 \cdot 0 + 1 \cdot (-1) + 1 \cdot 2 + 2 \cdot (-1) + 3 \cdot 0)$$

1D Convolution

- 1D convolution example:

- Signal ‘x’:

0	1	1	2	3	5	8	13
---	---	---	---	---	---	---	----

- Filter ‘w’:

0	-1	2	-1	0
---	----	---	----	---

- Convolution ‘z’:

		-1	0				
--	--	----	---	--	--	--	--

1D Convolution

- 1D convolution example:

- Signal ‘x’:

0	1	1	2	3	5	8	13
---	---	---	---	---	---	---	----

- Filter ‘w’:

0	-1	2	-1	0
---	----	---	----	---

- Convolution ‘z’:

		-1	0	-1			
--	--	----	---	----	--	--	--

1D Convolution

- 1D convolution example:

- Signal ‘x’:

0	1	1	2	3	5	8	13
---	---	---	---	---	---	---	----

- Filter ‘w’:

0	-1	2	-1	0
---	----	---	----	---

- Convolution ‘z’:

		-1	0	-1	-1		
--	--	----	---	----	----	--	--

1D Convolution Examples

- Examples:

- “Identity”

$$\hookrightarrow w = [0 \ 1 \ 0]$$

Let $x = [0 \ 1 \ 1 \ 2 \ 3 \ 5 \ 8 \ 13]$

$$z = [0 \ 1 \ 1 \ 2 \ 3 \ 5 \ 8 \ 13]$$

$0 \cdot x_0 + 1 \cdot x_1 + 0 \cdot x_2$ $0 \cdot x_1 + 1 \cdot x_2 + 0 \cdot x_3$

- “Translation”

$$\hookrightarrow w = [0 \ 0 \ 1]$$

$$z = [1 \ 1 \ 2 \ 3 \ 5 \ 8 \ 13 \ ?]$$

$0 \cdot x_0 + 0 \cdot x_1 + 1 \cdot x_2$

1D Convolution Examples

- Examples:
 - “Identity”

$$\hookrightarrow w = [0 \ 1 \ 0]$$

Let $x = [0 \ 1 \ 1 \ 2 \ 3 \ 5 \ 8 \ 13]$

average

average

- “Local Average”

$$\hookrightarrow w = [\frac{1}{3} \ \frac{1}{3} \ \frac{1}{3}]$$

$$z = [? \ \frac{2}{3} \ \frac{1}{3} \ 2 \ \frac{3\frac{1}{3}}{3} \ \frac{5\frac{1}{3}}{3} \ \frac{8\frac{2}{3}}{3} ?]$$

Boundary Issue

- What can we do about the “?” at the edges?

If $x = [0 \ 1 \ 1 \ 2 \ 3 \ 5 \ 8 \ 13]$ and $w = [\frac{1}{3} \ \frac{1}{3} \ \frac{1}{3}]$ then $z = [? \ \frac{2}{3} \ 1\frac{1}{3} \ 2 \ 3\frac{1}{3} \ 5\frac{1}{3} \ 8\frac{2}{3} \ ?]$

- Can assign values past the boundaries:

- “Zero”: $x = [0 \ 0 \ 0 \ 0 \ 1 \ 1 \ 1 \ 2 \ 3 \ 5 \ 8 \ 13] \ 0 \ 0 \ 0$

- “Replicate”: $x = [0 \ 0 \ 0 \ 0 \ 1 \ 1 \ 1 \ 2 \ 3 \ 5 \ 8 \ 13] \ 13 \ 13 \ 13$

- “Mirror”: $x = [2 \ 1 \ 1 \ 0 \ 1 \ 1 \ 1 \ 2 \ 3 \ 5 \ 8 \ 13] \ 8 \ 5 \ 3$

- Or just ignore the “?” values and return a shorter vector:

$$z = [\frac{2}{3} \ 1\frac{1}{3} \ 2 \ 3\frac{1}{3} \ 5\frac{1}{3} \ 8\frac{2}{3}]$$

bonus!

Formal Convolution Definition

- We've defined the convolution as:

$$z_i = \sum_{j=-m}^m w_j x_{i+j}$$

- In other classes you may see it defined as:

$$z_i = \sum_{j=-m}^m w_j x_{i-j}$$

(reverses 'w')

$$z_i = \sum_{j=-\infty}^{\infty} w_j x_{i-j} d_j$$



(assumes signal + filter are continuous)

- For simplicity we're skipping the "reverse" step, and assuming 'w' and 'x' are sampled at discrete points (not functions).
- But keep this mind if you read about convolutions elsewhere.

1D Convolution Examples

- Translation convolution shift signal:
 - “What is my neighbour’s value?”



$$W = [1 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0 \ 0]$$



1D Convolution Examples



- Averaging convolution (“is signal generally high in this region?”)
 - Less sensitive to noise (or spikes) than raw signal.

$$w = \left[\frac{1}{9}, \frac{1}{9}, \frac{1}{9}, \frac{1}{9}, \frac{1}{9}, \frac{1}{9}, \frac{1}{9}, \frac{1}{9}, \frac{1}{9} \right]$$



1D Convolution Examples



- **Gaussian convolution (“blurring”):** $w_i \propto \exp\left(-\frac{i^2}{2\sigma^2}\right)$
 - Compared to averaging it's more smooth and maintains peaks better.
- $w = [0.0001 \ 0.0644 \ 0.0540 \ 0.1420 \ 0.3989 \ 0.2420 \ 0.0540 \ 0.0044 \ 0.0001]$
 $(o = 1, m = 4)$



1D Convolution Examples

- **Sharpen** convolution enhances peaks.
 - An “average” that places **negative weights** on the surrounding pixels.




$$w = [-1 \quad 3 \quad -1]$$



1D Convolution Examples

- Centered difference convolution approximates first derivative:
 - Positive means change from low to high (negative means high to low).


$$w = [-1 \quad 0 \quad 1]$$



Digression: Derivatives and Integrals


- Numerical derivative approximations can be viewed as filters:

- Centered difference: $[-1, 0, 1]$
(like `check_correctness` in the homework code)



- Numerical integration approximations can be viewed as filters:

- “Simpson’s” rule: $[1/6, 4/6, 1/6]$ (a bit like Gaussian filter).





- Derivative filters add to 0, integration filters add to 1

- For constant function, derivative should be 0 and average = constant.

1D Convolution Examples

- Laplacian convolution approximates second derivative:
 - “Sum to zero” filters “respond” if input vector looks like the filter

$$w = [-1 \quad 2 \quad -1]$$




Laplacian of Gaussian Filter

- Laplacian of Gaussian is a smoothed 2nd-derivative approximation:



$$w_i = \left(1 - \frac{i^2}{2\sigma^2}\right) \exp\left(-\frac{i^2}{2\sigma^2}\right)$$

(then subtract mean)



$$w = [-0.1416 \ -0.1781 \ -0.2746 \ 0.1640 \ 0.8607 \ 0.1640 \ -0.2746 \ -0.1781 \ -0.1416]$$

$(\sigma^2=1, m=4)$



Images and Higher-Order Convolution

- **2D convolution:**
 - Signal ‘x’ is the pixel intensities in an ‘n’ by ‘n’ image.
 - Filter ‘w’ is the pixel intensities in a ‘2m+1’ by ‘2m+1’ image.
- The **2D convolution** is given by:

$$z[i_1, i_2] = \sum_{j_1=-m}^m \sum_{j_2=-m}^m w[j_1, j_2] \times [i_1 + j_1, i_2 + j_2]$$

- **3D and higher-order convolutions** are defined similarly.

$$z[i_1, i_2, i_3] = \sum_{j_1=-m}^m \sum_{j_2=-m}^m \sum_{j_3=-m}^m w[j_1, j_2, j_3] \times [i_1 + j_1, i_2 + j_2, i_3 + j_3]$$






Image Convolution Examples

x



Identity convolution:
(zeroes with a '1' at $w_{0,0}$)

$$w \quad *$$

A small square black box representing a convolution kernel. It has a single white dot in the center. To its left is a multiplication symbol (*), and to its right is an equals sign (=).

multiply element-wise
and add up result to get





Image Convolution Examples

x



Identity convolution:
(zeroes with a '1' at $w_{0,0}$)

$$w \quad *$$

=

↓
multiply element-wise
and add up result to get



Image Convolution Examples



Translation Convolution:

$$\begin{matrix} * \\ \text{Boundary: "zero"} \end{matrix} \quad \begin{matrix} \text{Translation Convolution:} \\ \text{Boundary: "zero"} \end{matrix} = \begin{matrix} \text{Resulting Image} \end{matrix}$$

A diagram illustrating a convolution operation. On the left, there is a small black square with a red circle at its top-left corner, representing a receptive field boundary. To its left is a symbol resembling a star with a diagonal line through it. To its right is an equals sign. To the right of the equals sign is a larger black and white photograph of the same Shiba Inu dog, identical to the one on the far left. Below the first black square is the text "Boundary: 'zero'". Below the second black square is the text "Resulting Image".



Image Convolution Examples



Translation Convolution:

$$\begin{matrix} * \\ \text{---} \\ \ast \end{matrix} \quad \begin{matrix} \text{---} \\ \text{---} \\ \text{---} \end{matrix} = \begin{matrix} \text{---} \\ \text{---} \\ \text{---} \end{matrix}$$

Boundary: "replicate"




Image Convolution Examples



Translation Convolution:

$$\begin{matrix} * \\ \text{Input Image} \end{matrix} \quad = \quad \text{Output Image}$$

Boundary: "mirror"



Image Convolution Examples



Translation Convolution:

$$\begin{matrix} * & \begin{matrix} \textcircled{0} \\ \text{---} \\ \text{---} \end{matrix} & = & \begin{matrix} \text{---} \\ \text{---} \\ \text{---} \end{matrix} \end{matrix}$$

Boundary: "ignore"



Image Convolution Examples



Average convolution:

$$\ast \frac{1}{5} \begin{bmatrix} 1 & 1 & 1 & \cdots & 1 \\ 1 & 1 & 1 & \searrow & 1 \\ 1 & 1 & 1 & \cdots & 1 \\ 1 & 1 & 1 & \cdots & 1 \end{bmatrix} =$$





Image Convolution Examples



Gaussian Convolution:

*



=

blurs image to represent
average
(smoothing)





Image Convolution Examples



Gaussian Convolution:

*



=

(smaller variance)

blurs image to represent
average
(smoothing)

A handwritten mathematical expression illustrating a convolution operation. It shows a large input image followed by a multiplication sign (*), then a small square kernel containing a single white pixel at its center, followed by an equals sign (=). Below the kernel is the text "(smaller variance)". Below the entire expression is the explanatory text "blurs image to represent average (smoothing)".

Image Convolution Examples



Laplacian of Gaussian

$$\begin{matrix} * & & \\ & \cdot & \\ & & \end{matrix} =$$

"How much does it look
like a black dot
surrounded by white?"




Image Convolution Examples



Laplacian of Gaussian

*



=

(larger variance)

Similar preprocessing may be
done in basal ganglia and LGN.



Image Convolution Examples



"Emboss" filter:

$$\ast \begin{bmatrix} -2 & -1 & 0 \\ -1 & 0 & 1 \\ 0 & 1 & 2 \end{bmatrix} =$$

Many Photoshop effects
are just convolutions.



Image Convolution Examples



Gabor filter
(Gaussian multiplied by
Sine or cosine)

$$\ast \quad \begin{matrix} & \\ & \text{A small, localized bright spot on a dark background} \\ & \end{matrix} =$$

$$\begin{matrix} & \text{A square containing concentric circles} \\ \text{Gaussian} & \ast \\ & \end{matrix} \quad \begin{matrix} & \text{A square containing several parallel wavy lines} \\ \text{Parallel Sine functions} & \end{matrix}$$





Image Convolution Examples




Gabor filter
(Gaussian multiplied by
Sine or cosine)

*



=



Different orientations of
the sine/cosine let us
detect changes with different
 orientations.


→ 2d derivatives have a direction.

Image Convolution Examples



Gabor filter
(Gaussian multiplied by
Sine or cosine)

*



=

(smaller variance)





Image Convolution Examples



Gabor filter
(Gaussian multiplied by
Sine or cosine)

*



=

(smaller variance)

Vertical orientation

- Can obtain other orientations by rotating.



- May be similar to effect of V1 "simple cells."




Image Convolution Examples



Max absolute value
between horizontal and
vertical Gabor:



maximum
absolute
value




"Horizontal/vertical edge detector"

3D Convolution



Represent
as RGB




Can apply 3D
convolutions

3D Convolution




Gaussian filter




3D Convolution




Gaussian filter
(higher variance on
green channel)



3D Convolution



Sharpen the blue channel.



3D Convolution




Gabor filter on
each channel.




Convolutions as Features

- Classic vision methods use **convolutions as features**:
 - Usually have different types/variances/orientations.
 - Can take maxes across locations/orientations/scales.
- Notable convolutions:
 - Gaussian** (blurring/averaging).
 - Laplace of Gaussian** (second-derivative).
 - Gabor filters** (directional first- or higher-derivative).



Filter Banks

- To characterize context, we used to use **filter banks** like “MR8”:
 - 1 Gaussian filter, 1 Laplacian of Gaussian filter.
 - 6 max(abs(Gabor)) filters:
 - 3 scales of sine/cosine (maxed over 6 orientations).



- **Convolutional neural networks** (next time!) are replacing filter banks.

Summary

- **Convolutions** are flexible class of signal/image transformations.
 - Can approximate directional derivatives and integrals at different scales.
 - **Max(convolutions)** can yield features invariant to some transformations.
- **Filter banks:**
 - Make features for a vision problem by takin a bunch of convolutions.
- Next time:
 - Combining this with deep learning.

bonus!

Global and Local Features for Domain Adaptation

- Suppose you want to solve a classification task, where you have very little labeled data from your domain.
- But you have access to a huge dataset with the same labels, from a different domain.
- Example:
 - You want to label POS tags in medical articles, and pay a few \$\$\$ to label some.
 - You have access the thousands of examples of Wall Street Journal POS labels.
- **Domain adaptation:** using data from different domain to help.

bonus!

Global and Local Features for Domain Adaptation

- “Frustratingly easy domain adaptation”:
 - Use “global” features across the domains, and “local” features for each domain.
 - “Global” features let you learn patterns that occur across domains.
 - Leads to sensible predictions for new domains without any data.
 - “Local” features let you learn patterns specific to each domain.
 - Improves accuracy on particular domains where you have more data.
 - For linear classifiers this would look like:


$$\hat{y}_i = \text{sign}(w_g^\top x_{ig} + w_d^\top x_{id})$$

features used across domains features/weights specific to domain.

bonus!

Image Coordinates

- Should we use the image coordinates?
 - E.g., the pixel is at location (124, 78) in the image.




- Considerations:
 - Is the interpretation different in different areas of the image?
 - Are you using a linear model?
 - Would “distance to center” be more logical?
 - Do you have enough data to learn about all areas of the image?

bonus!

Alignment-Based Features

- The position in the image is important in brain tumour application.
 - But we didn't have much data, so **coordinates didn't make sense**.
- We aligned the images with a “template image”.



(Look different because
we're showing middle slice
and alignment is in 3D.)


bonus!

Alignment-Based Features

- The position in the image is important in brain tumour application.
 - But we didn't have much data, so **coordinates didn't make sense**.
- We aligned the images with a “template image”.
 - Allowed “alignment-based” features:

Probability of
gray matter at
this pixel among
tons of people aligned
with template.


Probability of
being brain pixel.



Motivation: Automatic Brain Tumor Segmentation

bonus!

- Final features for brain tumour segmentation:
 - Gaussian convolution of original/template/priors/symmetry, Laplacian of Gaussian on original.
 - All with 3 variances.
 - Max(Gabor) with sine and cosine on orginal (3 variances).



bonus!


Motivation: Automatic Brain Tumour Segmentation

- Logistic regression and SVMs among best methods.

- When using these 72 features from last slide.
 - If you used all features I came up with, it overfit.

- Possible solutions to overfitting:

- Forward selection was too slow.
 - Just one image gives 8 million training examples.
 - I did manual feature selection (“guess and check”).
 - L2-regularization with all features also worked.
 - But this is slow at test time.
 - L1-regularization gives best of regularization and feature selection.



bonus!

FFT implementation of convolution

- Convolutions can be implemented using fast Fourier transform:
 - Take FFT of image and filter, multiply elementwise, and take inverse FFT.
- It has faster asymptotic running time but there are some catches:
 - You need to be using periodic boundary conditions for the convolution.
 - Constants matter: it may not be faster in practice.
 - Especially compared to using GPUs to do the convolution in hardware.
 - The gains are largest for larger filters (compared to the image size).

bonus!

SIFT Features

- Scale-invariant feature transform (SIFT):
 - Features used for object detection (“is particular object in the image”?)
 - Designed to detect unique visual features of objects at multiple scales.
 - Proven useful for a variety of object detection tasks.

