

CPSC 340: Machine Learning and Data Mining

Convolutional Neural Networks

Fall 2021

Admin

- Lectures this week
 - Today: CNNs and a little more general deep learning
 - Wed: more CNNs, mostly fun examples
 - Fri: bonus lectures
 - 2pm: CNN details – how to actually do things in code
 - 4pm: generative adversarial nets – using CNNs to make fake people
- A6 updates/clarifications
 - See <https://piazza.com/class/ksums2w1qd91se?cid=639>

Recent Lectures: Deep Learning

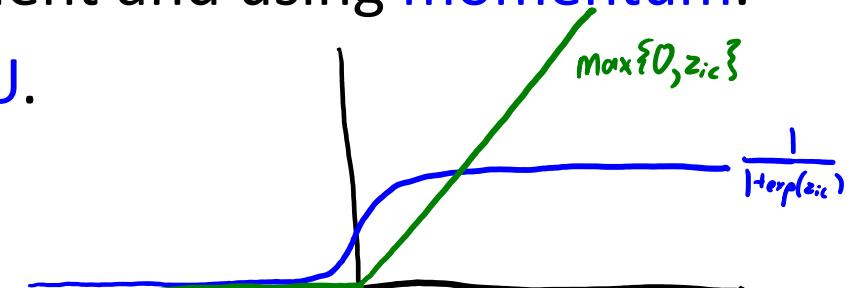
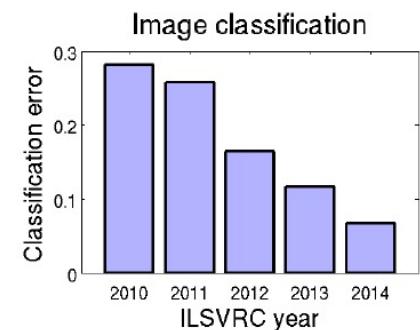
- We've been discussing **neural network / deep learning** models:

$$\hat{y}_i = v^T h(W^{(m)} h(W^{(m-1)} h(\dots \dots W^{(2)} h(W^{(1)} x_i)) \dots))$$

- We discussed unprecedented vision/speech performance.

- We discussed methods to make SGD work better:

- Parameter initialization and data transformations.
- Setting the step size(s) in stochastic gradient and using momentum.
- Alternative non-linear functions like ReLU.



“Residual” Networks (ResNets)

- Impactful recent idea is residual networks (**ResNets**):

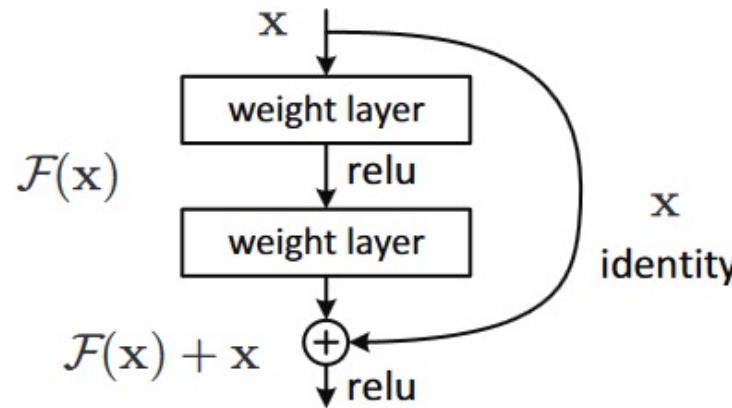


Figure 2. Residual learning: a building block.

- You can **take previous (non-transformed) layer as input** to current layer.
 - Also called “skip connections” or “highway networks”.
- **Non-linear part of the network only needs to model residuals.**
 - Non-linear parts are just “pushing up or down” a linear model in various places.
- This was a key idea behind first methods that used 100+ layers.
 - Evidence that biological networks have skip connections like this.

bonus!

DenseNet

- More recent variation is “DenseNets”:
 - Each layer can see all the values from many previous layers.
 - Gets rid of vanishing gradients.
 - May get same performance with fewer parameters/layers.

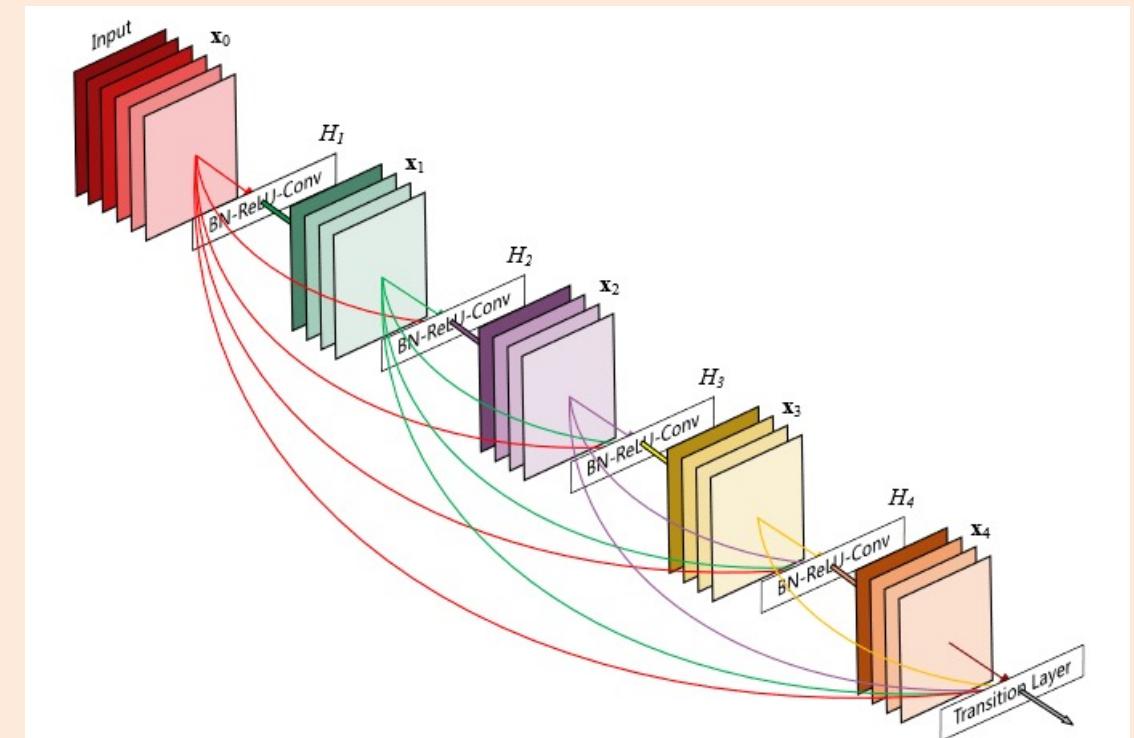


Figure 1: A 5-layer dense block with a growth rate of $k = 4$. Each layer takes all preceding feature-maps as input.

Deep Learning and the Fundamental Trade-Off

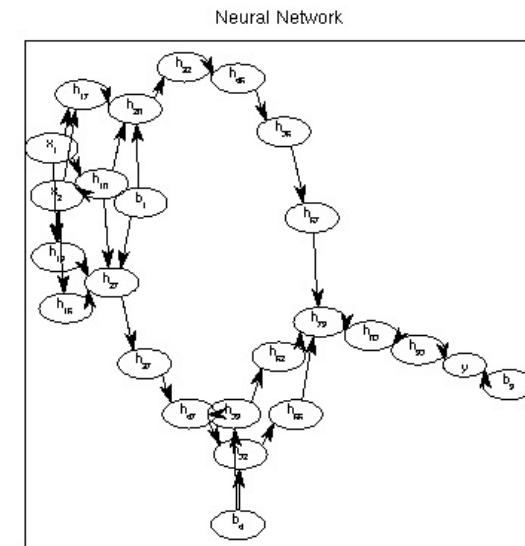
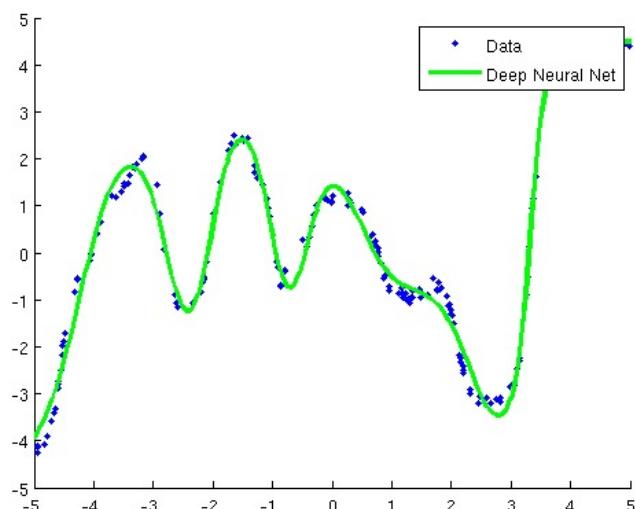
- Neural networks are subject to the fundamental trade-off:
 - With increasing depth, training error of global optima decreases.
 - With increasing depth, training error may poorly approximate test error.
- We want deep networks to model highly non-linear data.
 - But increasing the depth can lead to **overfitting**.
- How could GoogLeNet use 22 layers?
 - Many forms of **regularization** and keeping model complexity under control.
 - Unlike linear models, typically use **multiple types of regularization**.

Standard Regularization

- Traditionally, we've added our usual L2-regularizers:

$$f(v, W^{(3)}, W^{(2)}, W^{(1)}) = \frac{1}{2} \sum_{i=1}^n (v^\top h(W^{(3)} h(W^{(2)} h(W^{(1)} x_i))) - y_i)^2 + \frac{\lambda_4}{2} \|v\|^2 + \frac{\lambda_3}{2} \|W^{(3)}\|_F^2 + \frac{\lambda_2}{2} \|W^{(2)}\|_F^2 + \frac{\lambda_1}{2} \|W^{(1)}\|_F^2$$

- L2-regularization often called “weight decay” in this context.
 - Could also use L1-regularization: gives sparse network.



Standard Regularization

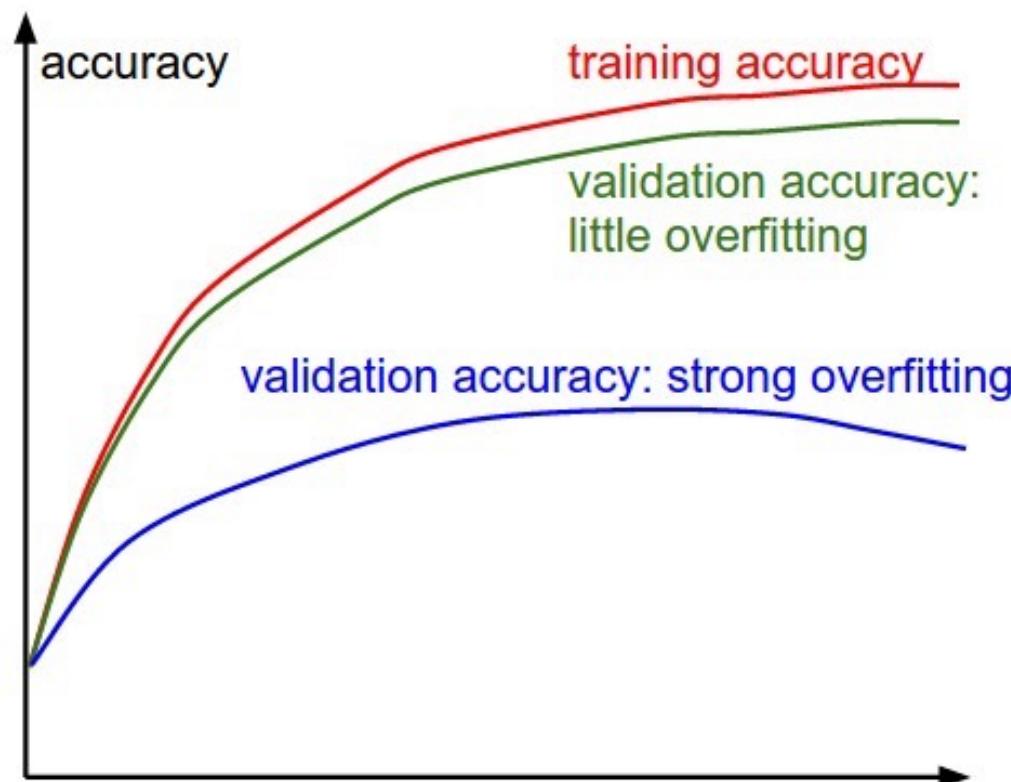
- Traditionally, we've added our usual L2-regularizers:

$$f(v, W^{(3)}, W^{(2)}, W^{(1)}) = \frac{1}{2} \sum_{i=1}^n (v^\top h(W^{(3)} h(W^{(2)} h(W^{(1)} x_i))) - y_i)^2 + \frac{\lambda_4}{2} \|v\|^2 + \frac{\lambda_3}{2} \|W^{(3)}\|_F^2 + \frac{\lambda_2}{2} \|W^{(2)}\|_F^2 + \frac{\lambda_1}{2} \|W^{(1)}\|_F^2$$

- L2-regularization often called “weight decay” in this context.
 - Adds λW to gradient, so (S)GD “decays” the weights ‘W’ at each step
 - Could also use L1-regularization: gives sparse network.
- Hyper-parameter optimization gets expensive:
 - Try to optimize validation error in terms of $\lambda_1, \lambda_2, \lambda_3, \lambda_4$.
 - In addition to step-size, number of layers, size of layers, initialization.
- Recent result:
 - Adding a regularizer in this way can create bad local optima.

Early Stopping

- Another common type of regularization is “early stopping”:
 - Monitor the validation error as we run stochastic gradient.
 - Stop the algorithm if validation error starts increasing.



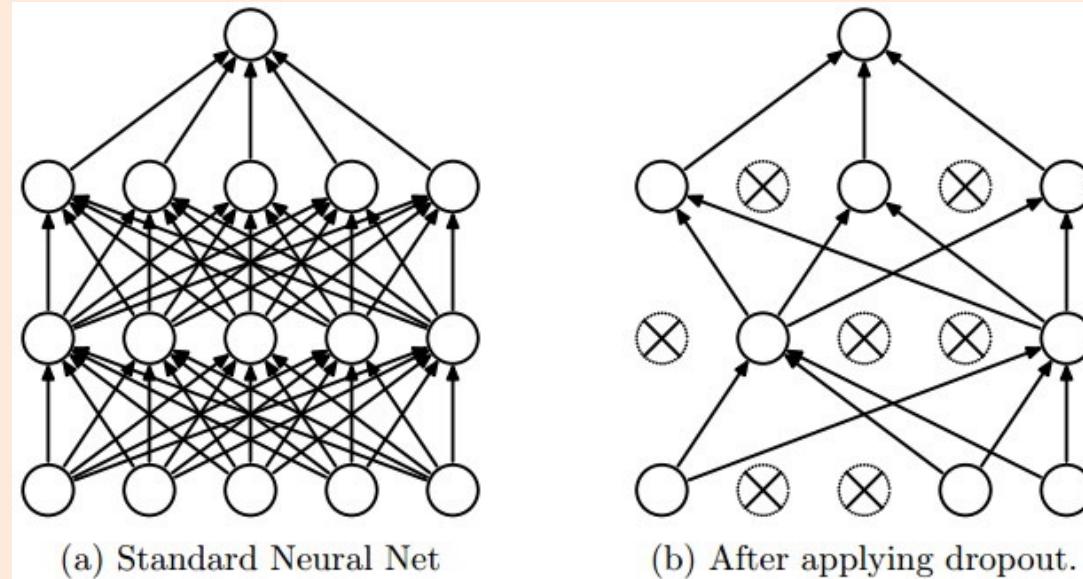
Unfortunately it might look more like

hopefully you don't stop here.

bonus!

Dropout

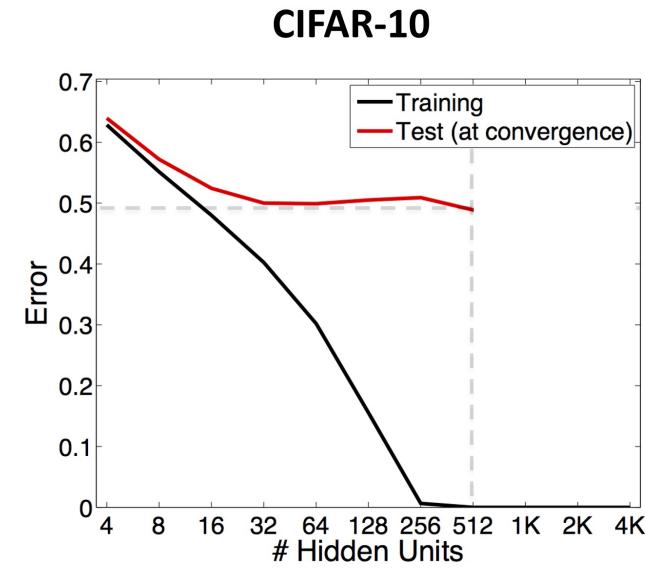
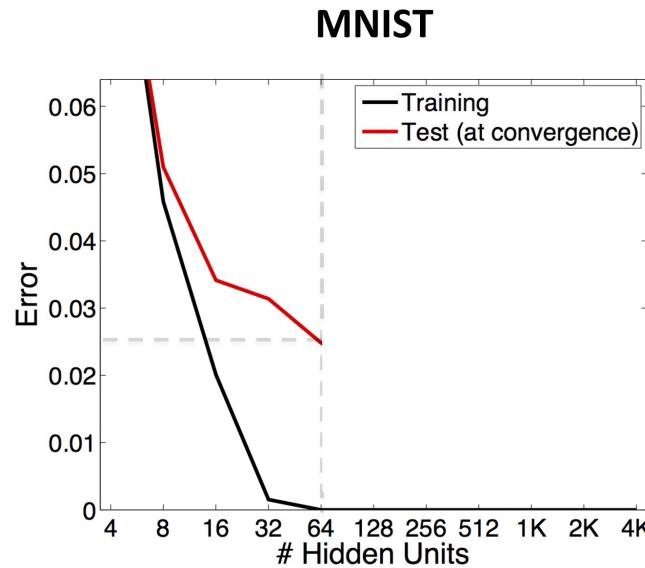
- **Dropout** is a more recent form of explicit regularization:
 - On each iteration, randomly set some x_i and z_i to zero (often use 50%).



- Adds invariance to missing inputs or latent factors
 - Encourages distributed representation rather than relying on specific z_i .
- Can be interpreted as an ensemble over networks with different parts missing.
- After a lot of early success, dropout is already kind of going out of fashion.

“Hidden” Regularization in Neural Networks

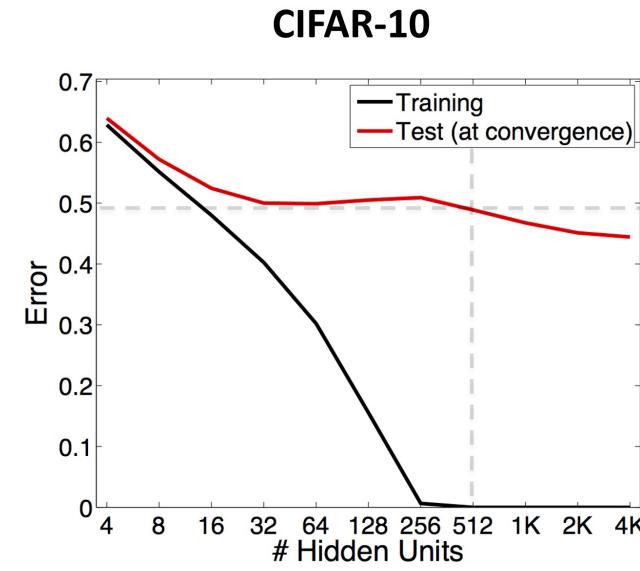
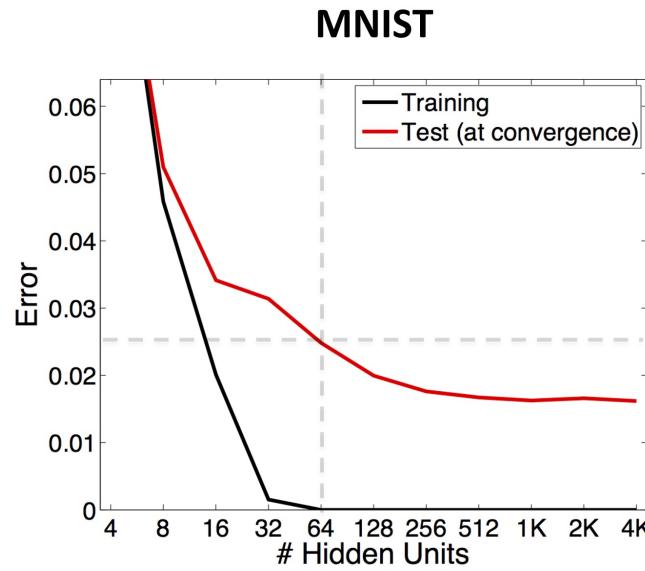
- Fitting single-layer neural network with SGD and no regularization:



- Training goes to 0 with enough units: we’re finding a global min.
- What should happen to training and test error for larger #hidden?

“Hidden” Regularization in Neural Networks

- Fitting single-layer neural network with SGD and no regularization:



- Test error continues to go down!?! Where is fundamental trade-off??
- There exist global mins with large #hidden units have test error = 1.
 - But among the global minima, SGD is somehow converging to “good” ones.

bonus!

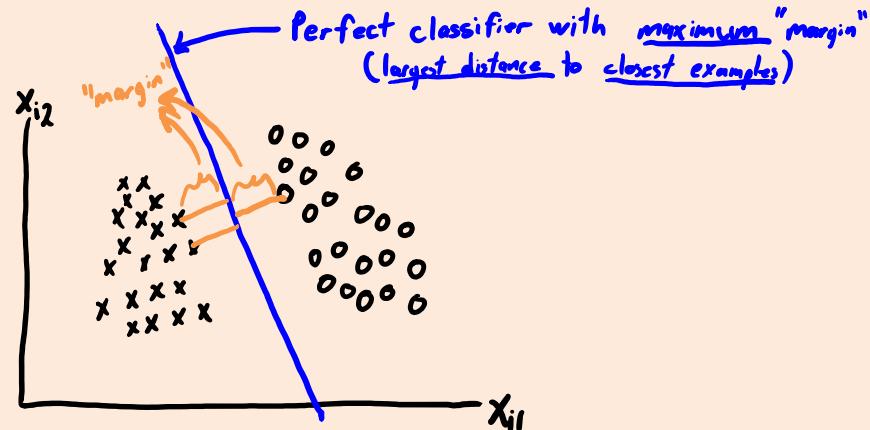
Implicit Regularization of SGD

- There is growing evidence that using SGD regularizes parameters.
 - We call this the “implicit regularization” of the optimization algorithm.
- Beyond empirical evidence, we know this happens in simpler cases.
- Example of implicit regularization:
 - Consider a least squares problem where there exists a ‘w’ where $X w = y$.
 - Residuals are all zero, we fit the data exactly.
 - You run [stochastic] gradient descent starting from $w=0$.
 - Converges to solution $X w = y$ that has the minimum L2-norm.
 - So using SGD is equivalent to L2-regularization here, but regularization is “implicit”.

bonus!

Implicit Regularization of SGD

- Example of implicit regularization:
 - Consider a **logistic regression** problem where data is linearly separable.
 - We can fit the data exactly.
 - You run gradient descent from any starting point.
 - Converges to **max-margin solution** of the problem.
 - So using gradient descent is equivalent to encouraging large margin.



- Similar result known for **boosting**.

(pause)

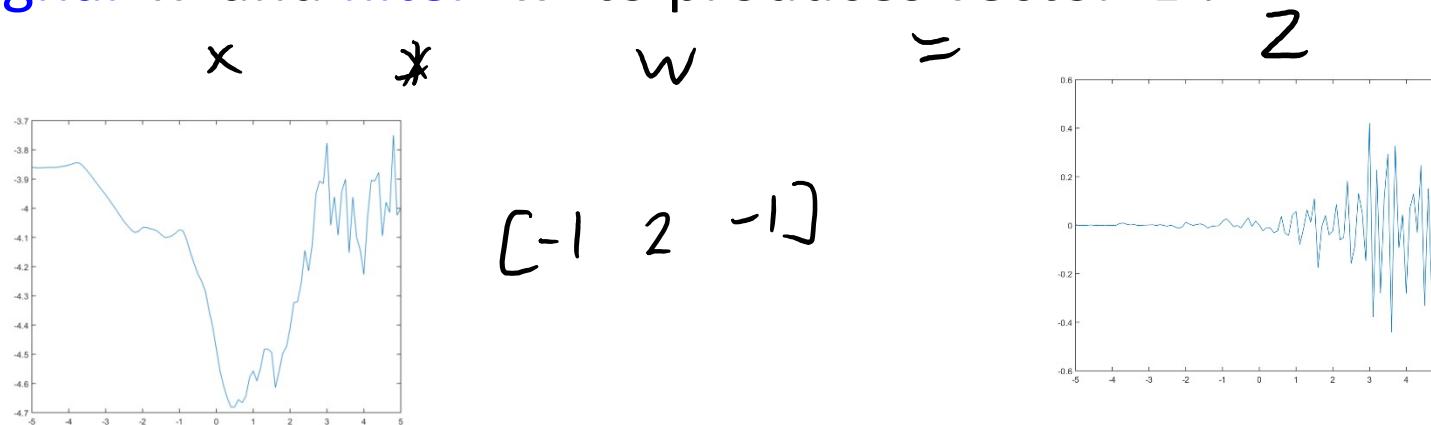
Deep Learning “Tricks of the Trade”

- We've discussed **heuristics to make deep learning work:**
 - Parameter initialization and data transformations.
 - Setting the **step size(s)** in stochastic gradient and using **momentum**.
 - **ResNets** and alternative non-linear functions like **ReLU**.
 - Different forms of regularization:
 - L2-regularization, early stopping, dropout, implicit regularization from SGD.
- These are often **still not enough** to get deep models working.
- Deep computer vision models are all **convolutional neural networks**:
 - The $W^{(m)}$ are **very sparse** and have **repeated parameters** (“tied weights”).
 - Drastically reduces number of parameters (speeds training, reduces overfitting).

1D Convolution as Matrix Multiplication

- **1D convolution:**

- Takes **signal** ‘x’ and **filter** ‘w’ to produces vector ‘z’:



- Can be written as a **matrix multiplication**:

$$W_x = \begin{bmatrix} 1 & -2 & 1 & 0 & 0 & 0 & \cdots & 0 & 0 & 0 & 0 \\ 0 & 1 & -2 & 1 & 0 & 0 & \cdots & 0 & 0 & 0 & 0 \\ 0 & 0 & 1 & -2 & 1 & 0 & \cdots & 0 & 0 & 0 & 0 \\ \vdots & & & & & & & & & & \\ 0 & 0 & 0 & 0 & 0 & 0 & \cdots & 0 & 1 & -2 & 1 \end{bmatrix} x = z$$

1D Convolution as Matrix Multiplication

- Each element of a convolution is an **inner product**:

$$\begin{aligned} z_i &= \sum_{j=-m}^m w_j x_{i+j} \\ &= w^T x_{(i-m:i+m)} \\ &= \tilde{w}^T x \quad \text{where } \tilde{w} = [0 \ 0 \ 0 \ \underbrace{w}_{\substack{\text{positions } i-m \text{ through } i+m}} \ 0 \ 0] \end{aligned}$$

- So **convolution is a matrix multiplication** (I'm ignoring boundaries):

$$z = \tilde{W}x \quad \text{where } \tilde{W} = \left[\begin{array}{cccccc} \overbrace{w} & & & & 0 & 0 & 0 \\ 0 & \overbrace{w} & & & 0 & 0 & \\ 0 & 0 & \overbrace{w} & & 0 & & \\ 0 & 0 & 0 & \overbrace{w} & & & \end{array} \right]$$

} matrix can be very sparse and only has $2m+1$ variables.

- The shorter 'w' is, the more sparse the matrix is.

2D Convolution as Matrix Multiplication

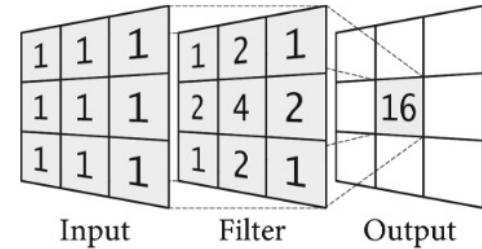
- **2D convolution:**

- Signal 'x', filter 'w', and output 'z' are now all **images/matrices**:

$$x * w = z$$



$$\begin{bmatrix} -2 & -1 & 0 \\ -1 & 0 & 1 \\ 0 & 1 & 2 \end{bmatrix}$$



- Vectorized 'z' can be written as a **matrix multiplication** with vectorized 'x':

$$W = \begin{bmatrix} -2 & -1 & 0 & 0 & 0 & \dots & 0 & -1 & 0 & 1 & 0 & 0 & \dots & 0 & 0 & 1 & 2 & 0 & 0 & \dots & 0 & 0 \\ 0 & -2 & -1 & 0 & 0 & \dots & 0 & 0 & -1 & 0 & 1 & 0 & 0 & \dots & 0 & 0 & 0 & 1 & 2 & 0 & \dots & 0 & 0 \\ \vdots & \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \ddots & \vdots & \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & 0 & 0 & 0 & \dots & 0 & 2 & -1 & 0 & -6 & 0 & 0 & \dots & 0 & -1 & 0 & 1 & 0 & \dots & 0 & 0 & \dots & 0 & 1 & 2 \\ 0 & 0 & 0 & 0 & 0 & \dots & 0 & 0 & 2 & -1 & 0 & -6 & 0 & 0 & \dots & 0 & -1 & 0 & 1 & 0 & \dots & 0 & 0 & \dots & 0 & 1 & 2 \\ 0 & 0 & 0 & 0 & 0 & \dots & 0 & 0 & 0 & 2 & -1 & 0 & -6 & 0 & 0 & \dots & 0 & -1 & 0 & 1 & 0 & \dots & 0 & 0 & \dots & 0 & 1 & 2 \\ \vdots & \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \ddots & \vdots & \vdots & \vdots & \vdots & \vdots & \ddots & \vdots & \vdots & \ddots & \vdots & \vdots \end{bmatrix}$$

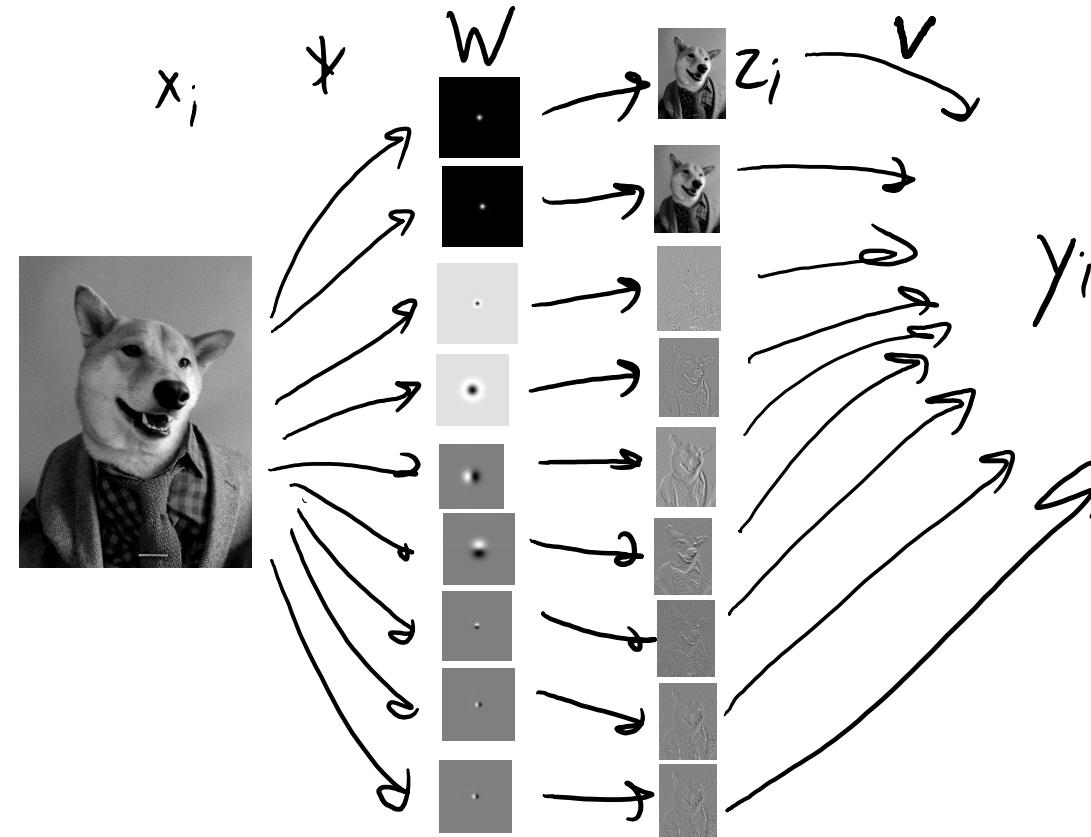
Motivation for Convolutional Neural Networks

- Consider training neural networks on 256 by 256 images.
 - This is 256 by 256 by 3 \approx 200,000 inputs.
- If first layer has $k=10,000$, then it has **about 2 billion parameters**.
 - We want to avoid this huge number (due to storage and overfitting).
- Key idea: make Wx_i act like several convolutions (to make it sparse):
 1. Each row of W only applies to part of x_i .
 2. Use the same parameters between rows.
- Forces most weights to be zero, reduces number of parameters.

$$w_1 = [0 \quad 0 \quad 0 \quad \dots \quad w \quad \dots \quad 0 \quad 0 \quad 0]$$
$$w_2 = [0 \quad \dots \quad w \quad \dots \quad 0 \quad 0 \quad 0 \quad 0]$$

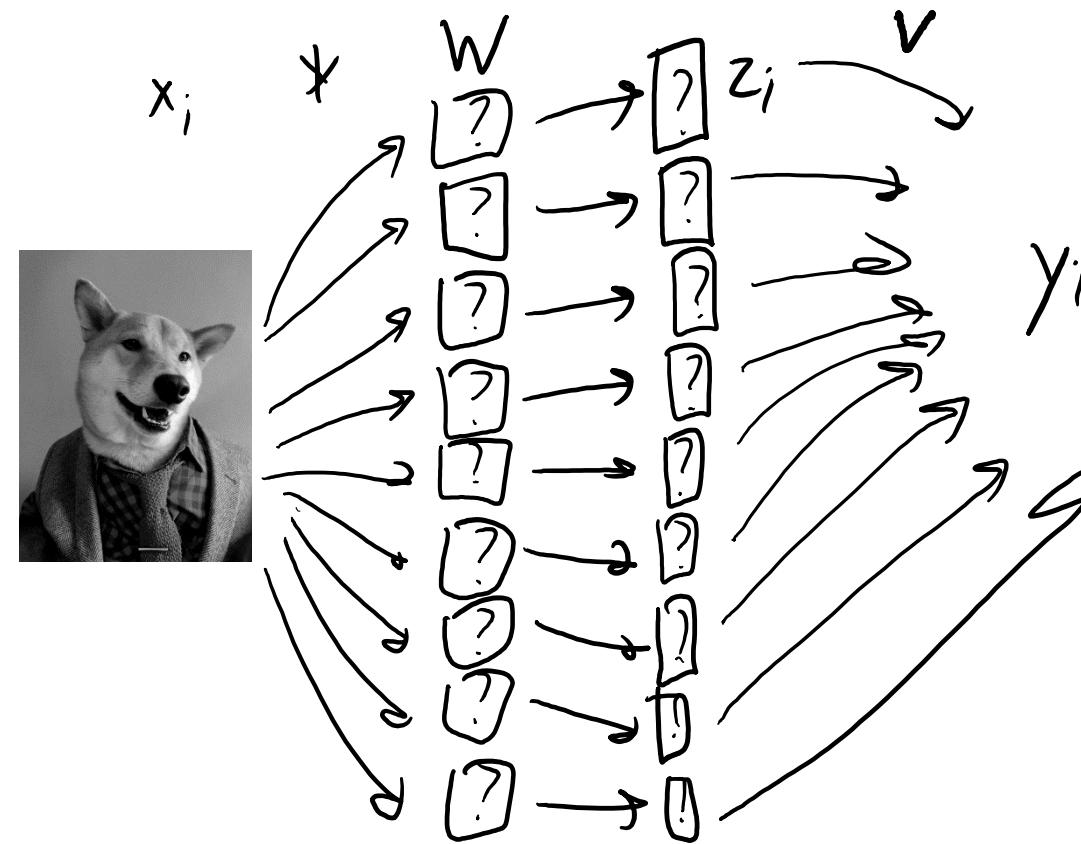
Motivation for Convolutional Neural Networks

- Classic vision methods uses **fixed convolutions** as features:
 - Usually have **different types/variances/orientations**.
 - Can do subsampling or take **maxes** across locations/orientations/scales.



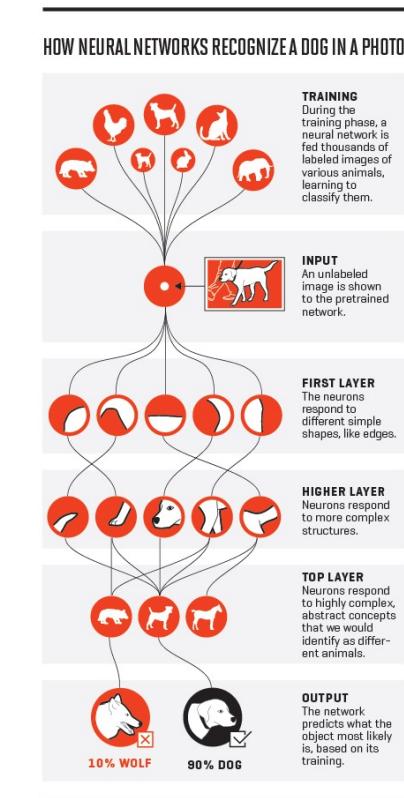
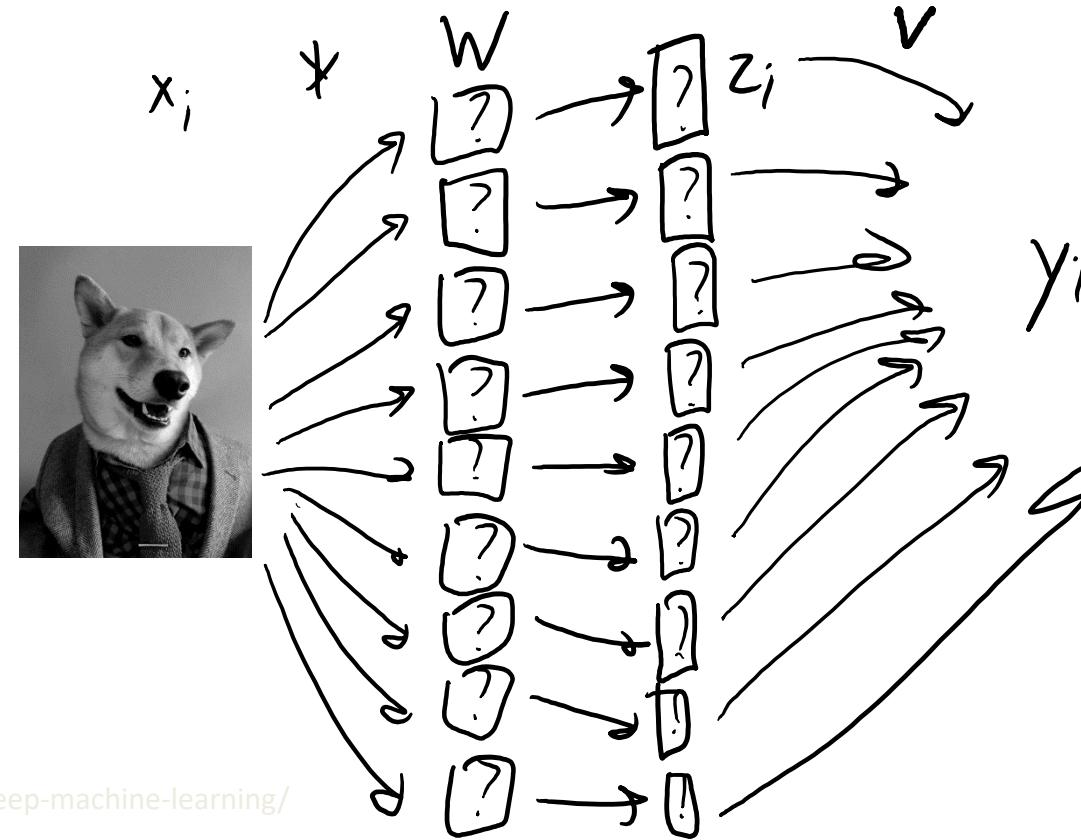
Motivation for Convolutional Neural Networks

- Convolutional neural networks learn the convolutions:
 - Learning ‘W’ and ‘v’ automatically chooses types/variances/orientations.
 - Don’t pick from fixed convolutions, but learn the elements of the filters.



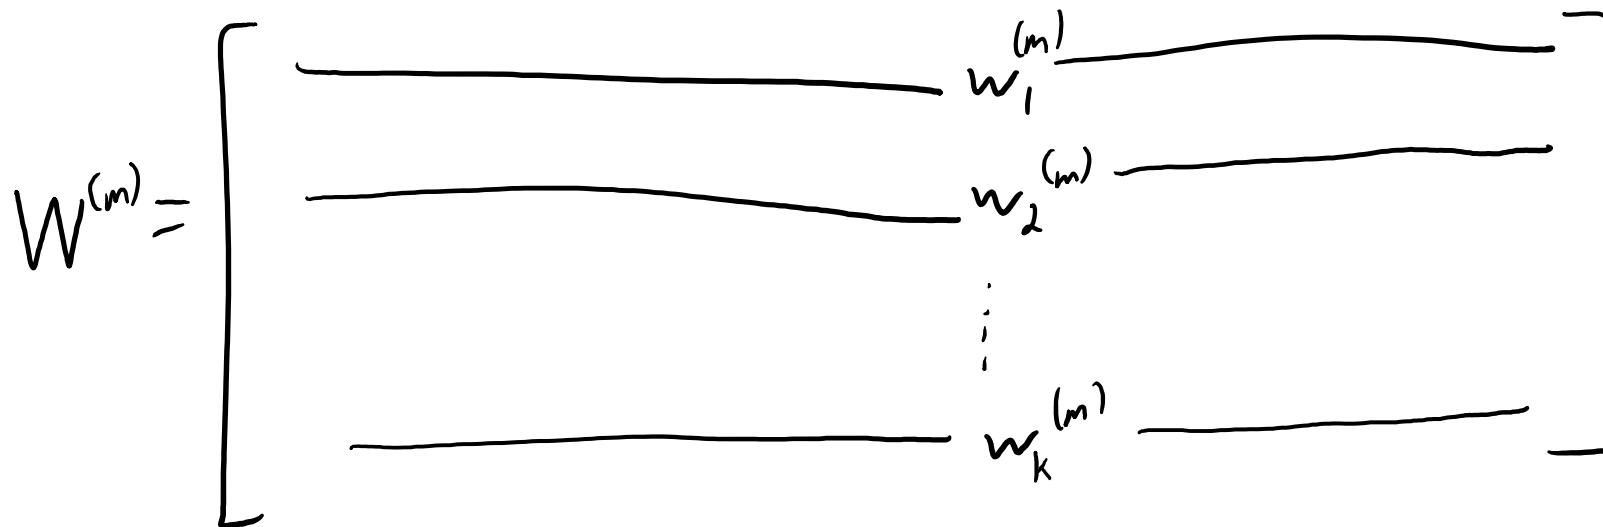
Motivation for Convolutional Neural Networks

- Convolutional neural networks learn the convolutions:
 - Learning ‘W’ and ‘v’ automatically chooses types/variances/orientations.
 - Can do multiple layers of convolution to get deep hierarchical features.



Convolutional Neural Networks

- Convolutional Neural Networks classically have 3 layer “types”:
 - Fully connected layer: usual neural network layer with unrestricted W.



Convolutional Neural Networks

- Convolutional Neural Networks classically have 3 layer “types”:
 - Fully connected layer: usual neural network layer with unrestricted W.
 - Convolutional layer: restrict W to act like several convolutions.

1D example

$$W^{(m)} = \left[\begin{array}{ccccccccc} & w_1^{(m)} & & & & & & & \\ 0 & 0 & 0 & & w_1^{(m)} & & & & \\ & & & & & & & & \\ 0 & 0 & 0 & 0 & 0 & 0 & & & \\ & & & & & & w_1^{(m)} & & \\ \hline & w_2^{(m)} & & & & & 0 & 0 & 0 \\ 0 & 0 & 0 & & w_2^{(m)} & & & & \\ & & & & & & & & \\ 0 & 0 & 0 & 0 & 0 & 0 & & & \\ & & & & & & w_2^{(m)} & & \\ \hline 0 & 0 & 0 & & 0 & 0 & 0 & 0 & 0 \end{array} \right]$$

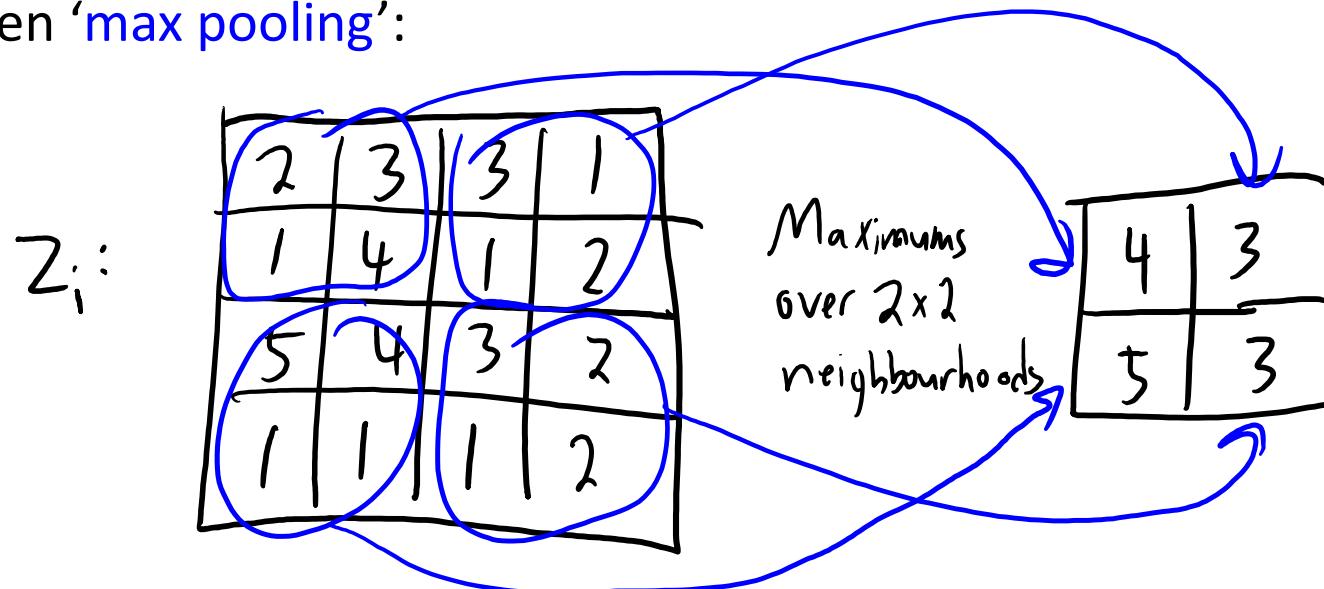
distance between centers of convolution is called "stride"

Same $w_i^{(m)}$ used across multiple rows.

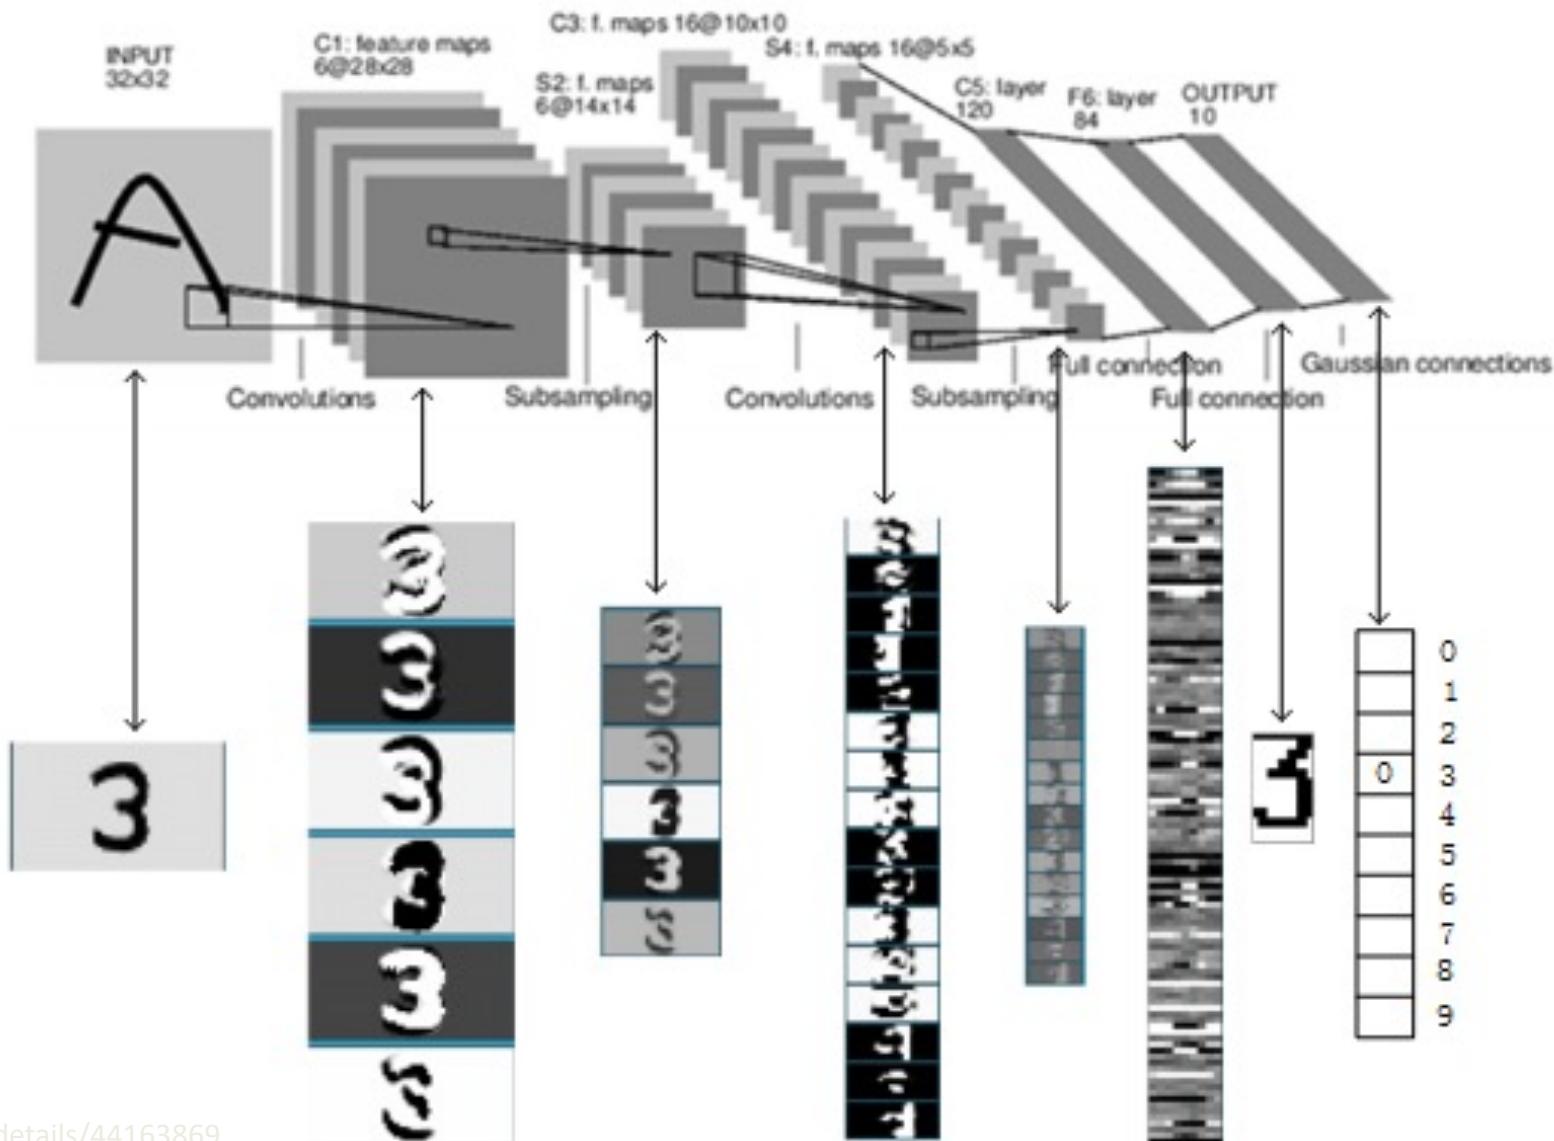
Sparse and small number of parameters

Convolutional Neural Networks

- Convolutional Neural Networks classically have 3 layer “types”:
 - Fully connected layer: usual neural network layer with unrestricted W.
 - Convolutional layer: restrict W to act like several convolutions.
 - Pooling layer: combine results of convolutions.
 - Can add some invariance or just make the number of parameters smaller.
 - Often ‘max pooling’:

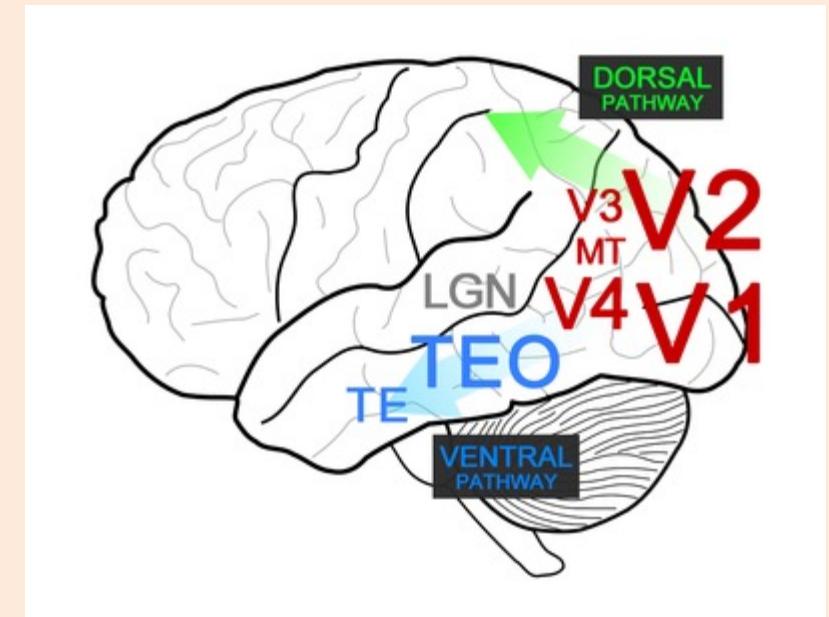
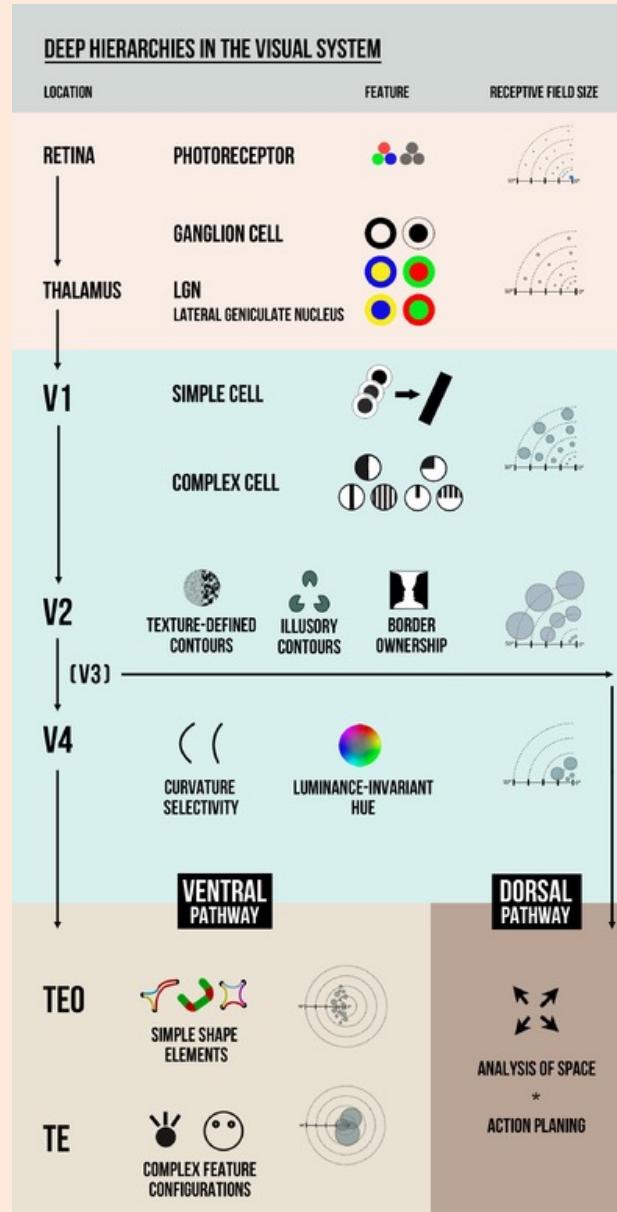
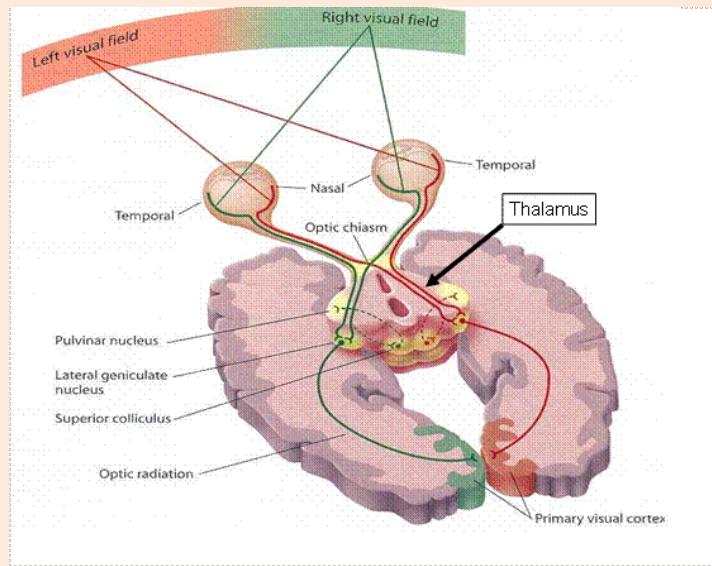


LeNet for Optical Character Recognition



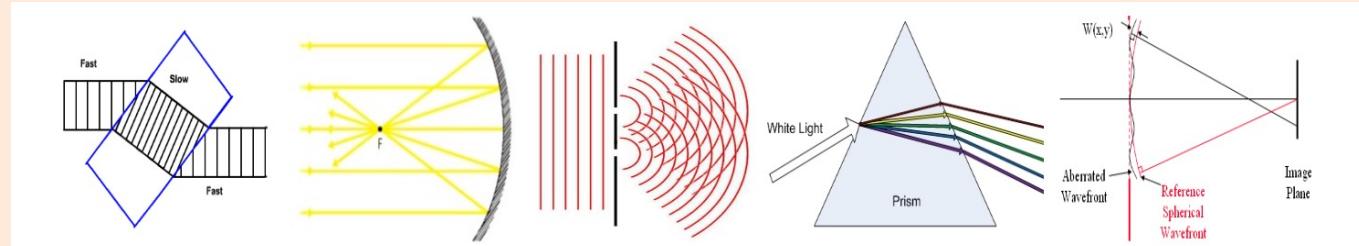
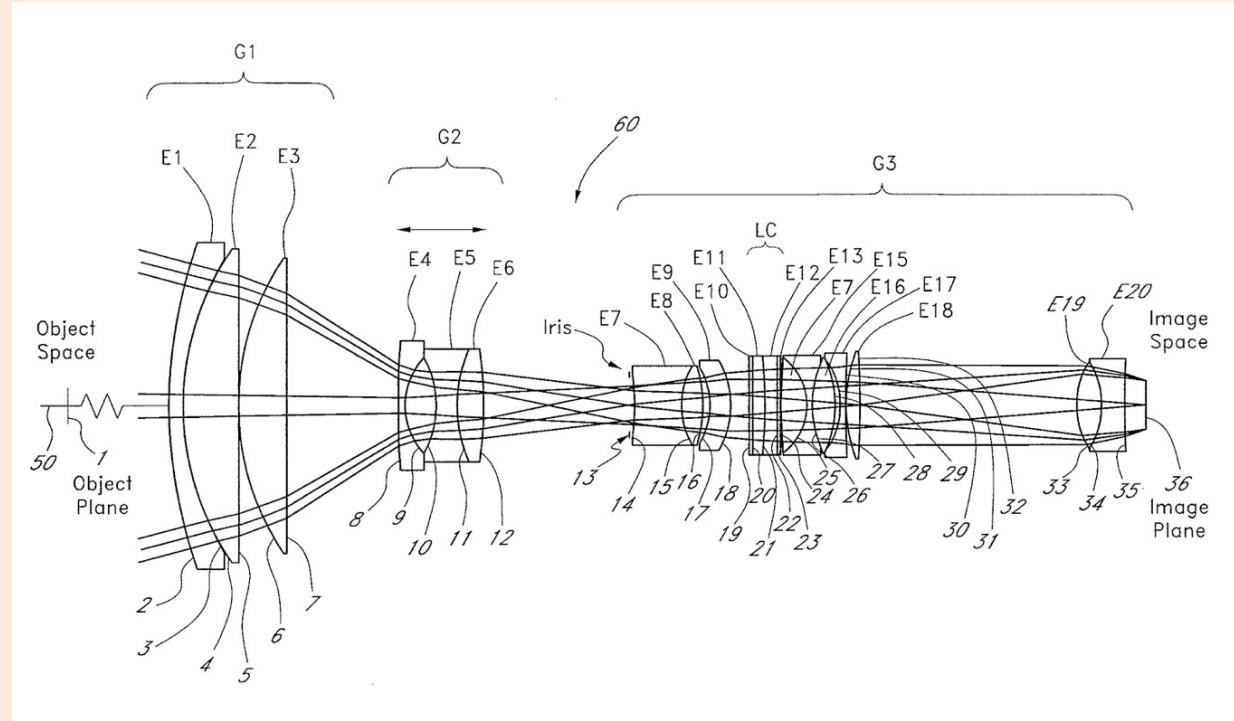
bonus!

Deep Hierarchies in the Visual System



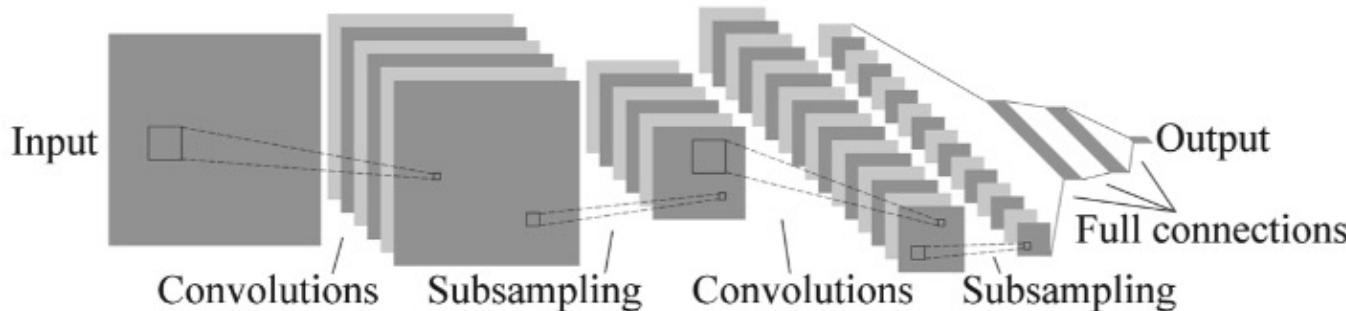
bonus!

Deep Hierarchies in Optics



Convolutional Neural Networks

- Classic convolutional neural network (LeNet):



- Visualizing the “activations” of the layers:

- <http://scs.ryerson.ca/~aharley/vis/conv>
- <http://cs231n.stanford.edu>



Summary

- **ResNets** include untransformed previous layers.
 - Network focuses non-linearity on residual, allows huge number of layers.
- **Regularization** is crucial to neural net performance:
 - L2-regularization, early stopping, dropout, implicit regularization of SGD.
- **Convolutional neural networks:**
 - Restrict $W^{(m)}$ matrices to represent sets of convolutions.
 - Often combined with max (pooling).
- Next time: modern convolutional neural networks and applications.
 - Image segmentation, depth estimation, image colorization, artistic style.