

CPSC 340 – Tutorial 5

Michael Liu
mfliu@students.cs.ubc.ca

Template from Lironne Kurzman

Some slides from Nam Hee Kim


University of British Columbia

November 1st, 2021

Agenda

- Multi-class classification
- Kernel Trick

Multi-class Classification



Motivation

- Our linear classifiers only work for binary classification
 - E.g. “important” vs “not important” emails
 - E.g. “cat” vs “no cat” in picture
- We would like to build a classifier that works for $k > 2$ classes
 - E.g. “cat” vs “dog” vs “bird” in picture

Notation Review

- X is an $n \times d$ matrix of **training examples**
- y is an $n \times 1$ vector of **training labels**
- x_i is a $d \times 1$ vector of **feature values** for example i
- y_i is a *scalar* in the range $[1, k]$, where there are k classes
- W is a $k \times d$ matrix of **weights** for each classifier $1 \leq c \leq k$
- w_c is a $d \times 1$ vector of **weights** for the classifier for class c
 - Can also be written as w_{y_i} where $y_i = c$

One-vs-all Linear Classification

- Idea: Independently train a linear classifier for each class c
 - Each classifier predicts how likely x_i is in class c
 - *E.g.* *Classifier 1: x_i is probably not a dog*
 Classifier 2: x_i is probably not a cat
 Classifier 3: x_i is probably a bird
 One-vs-all classifier: Predict x_i is a bird
 - At test time, use argmax to find “best class” for example \tilde{x}_i
- Problem: classifiers may operate with different scales
 - How do we compare $w_c^T x_i$ with $w_{c'}^T x_i$?

Multi-class SVM

- Goal: Every classifier outputs a value on the same scale
 - ☹️ Now we must train all k classifiers together
 - ☹️ W is no longer a simple vector
 - 😊 It performs better than one-vs-all linear classification
 - 😊 We only need to implement it once
- We want to make $w_{y_i}^T x_i$ larger than $w_c^T x_i$ for all $c \neq y_i$

Multi-class SVM Loss

$$\sum_{c \neq y_i} \max\{0, 1 - w_{y_i}^T x_i + w_c^T x_i\}$$

- This is the loss for a **single** example x_i
 - To get the loss over all examples, we can take the sum or average
- The 1 term avoid degeneracy (see binary SVM loss)
- The $-w_{y_i}^T x_i$ term **rewards** large values for the correct class
- The $w_c^T x_i$ term **penalizes** large values for incorrect classes


Kernel Trick

The background is split into two main sections. On the left, a large, light gray circle with a subtle gradient and thin white concentric lines is partially visible. On the right, a dark purple field is filled with a complex network of thin red lines that intersect to form a grid-like pattern. Scattered throughout this field are numerous small, bright yellow dots of varying sizes, some of which appear to be connected by the red lines, suggesting a network or data structure.

Motivation

- We've already learned about linear models
 - Easy to train
 - Easy to test
- We've already learned about change of basis
 - Training and testing are simple (just use linear model)
 - Creating Z matrix is easy, but can take a long time
 - $O(d^p)$ values for **each** example

Motivation as a Picture



<https://math.stackexchange.com/questions/353607/how-do-inner-product-space-determine-half-planes>

“Other” Normal Equations

- Assume L2-regularized least squares objective with basis Z :

$$f(v) = \frac{1}{2} \|Zv - Y\|^2 + \frac{\lambda}{2} \|v\|^2$$

- Normal equations to find minimum v :

$$v = (Z^T Z + \lambda I)^{-1} Z^T Y$$

- Other normal equations to find minimum v :

$$v = Z^T (Z Z^T + \lambda I)^{-1} Y$$

Making Predictions

- If we want to make predictions \hat{y} for test data \tilde{X} by forming \tilde{Z} using “other normal” equations:

$$\begin{aligned}\hat{y} &= \tilde{Z}v \\ &= \tilde{Z}\tilde{Z}^T(\tilde{Z}\tilde{Z}^T + \lambda I)^{-1}Y \\ &= \tilde{K}(K + \lambda I)^{-1}Y\end{aligned}$$

- Efficiently compute K and \tilde{K} even though forming Z and \tilde{Z} is intractable.

Kernel Functions

- Kernel function does not calculate Z explicitly. Returns similarity between transformed points Z_i, Z_j :

$$K(X_i, X_j) = Z_i^T Z_j$$

using only untransformed points X_i, X_j .

Example Kernels

- Linear Kernel:

$$K(X_i, X_j) = Z_i^T Z_j = X_i^T X_j$$

- Degree P polynomial Kernel:

$$K(X_i, X_j) = Z_i^T Z_j = (1 + X_i^T X_j)^P$$

- Gaussian RBF Kernel:

$$K(X_i, X_j) = Z_i^T Z_j = \exp\left(-\frac{\|X_i - X_j\|^2}{2\sigma^2}\right)$$